

Invertible Normalizing Flows

ECE57000: Artificial Intelligence

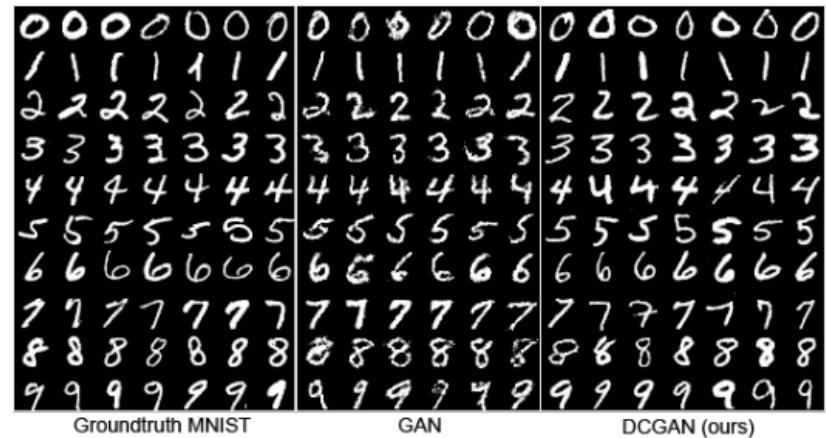
David I. Inouye

GAN Limitation:

Cannot compute density values

- ▶ Evaluation of GANs is challenging
 - ▶ Explicit density models could use test log likelihood
 - ▶ “I think this looks better than that”

- ▶ Inception-based scores



- ▶ Cannot use for classification or outlier detection

GAN Limitation: Challenging to train because of careful balance between discriminator and generator

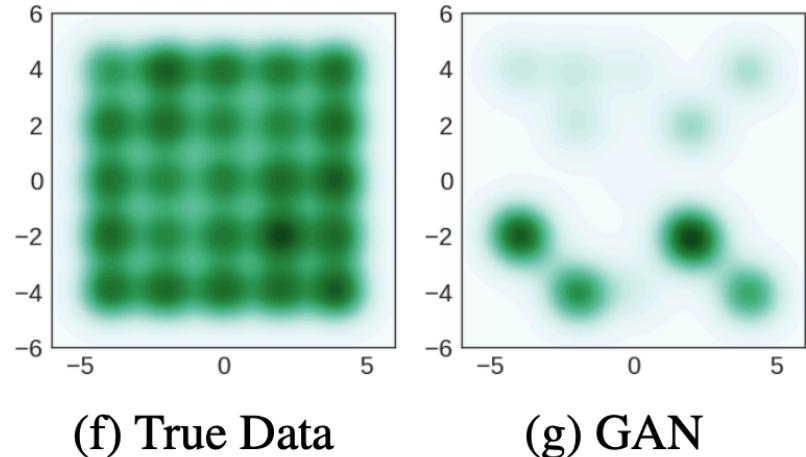
1. Assumptions on possible D and G
 1. Theory – All possible D and G
 2. Reality – Only functions defined by a neural network
 2. Assumptions on optimality
 1. Theory – Both optimizations are solved perfectly
 2. Reality – The inner maximization is only solved approximately, and this interacts with outer minimization
 3. Assumption on expectations
 1. Theory – Expectations over true distribution
 2. Reality – Empirical expectations over finite sample; for images, much of the high-dimensional space does not have samples
- **GANs can be very difficult/finicky to train**

Common problem with GANs: Mode collapse hinders diversity of samples

From: <https://developers.google.com/machine-learning/gan/problems>

- ▶ Wasserstein GANs
- ▶ Unrolled GANs
 - ▶ Trained on multiple discriminators simultaneously

Metz, L., Poole, B., Pfau, D., & Sohl-Dickstein, J. (2016). Unrolled generative adversarial networks. *arXiv preprint arXiv:1611.02163*.



<http://papers.nips.cc/paper/6923-veegan-reducing-mode-collapse-in-gans-using-implicit-variational-learning.pdf>

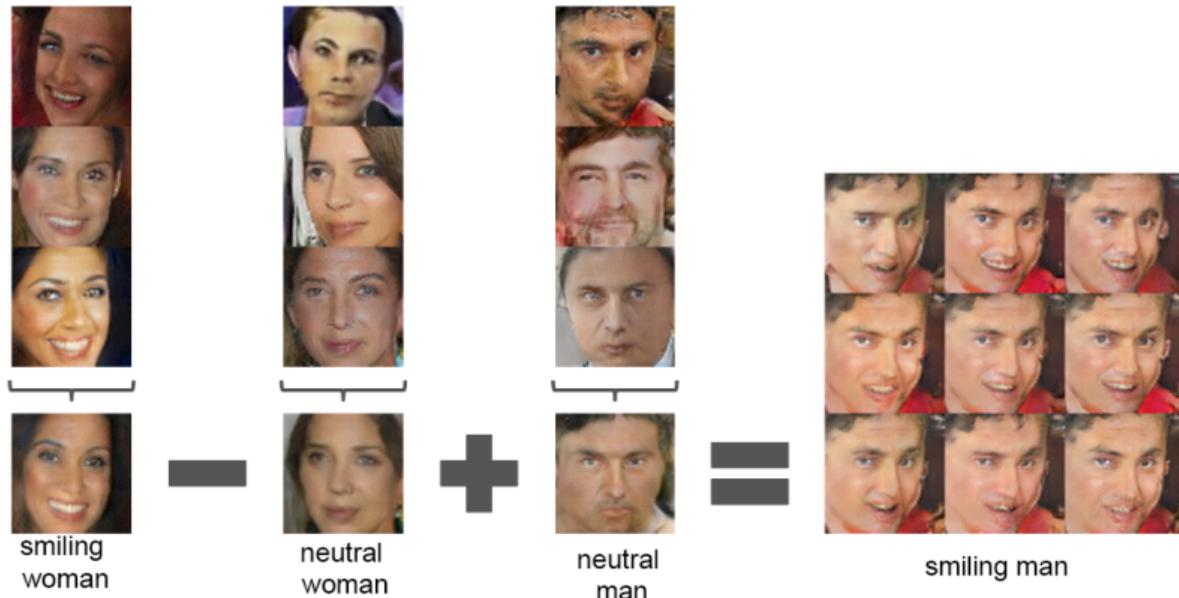


[https://software.intel.com/en-us/blogs/2017/08/21 mode-collapse-in-gans](https://software.intel.com/en-us/blogs/2017/08/21	mode-collapse-in-gans)

GAN Limitation: Cannot go from observed to latent space, i.e. $x \rightarrow z$ not possible/easy

- ▶ Cannot manipulate an observed image in latent space
 - ▶ Cannot do the following, $x \rightarrow z$, $z' = z + 3$, $z' \rightarrow x'$
 - ▶ Rather, must start from fake image based on random z

All fake
images->



Highly realistic random samples from powerful flow model (GLOW)



Figure 1: Synthetic celebrities sampled from our model; see Section 3 for architecture and method, and Section 5 for more results.

<https://papers.nips.cc/paper/8224-glow-generative-flow-with-invertible-1x1-convolutions.pdf>

Interpolation between real images using GLOW



Figure 5: Linear interpolation in latent space between real images.

<https://papers.nips.cc/paper/8224-glow-generative-flow-with-invertible-1x1-convolutions.pdf>

Transformations of real image along various features

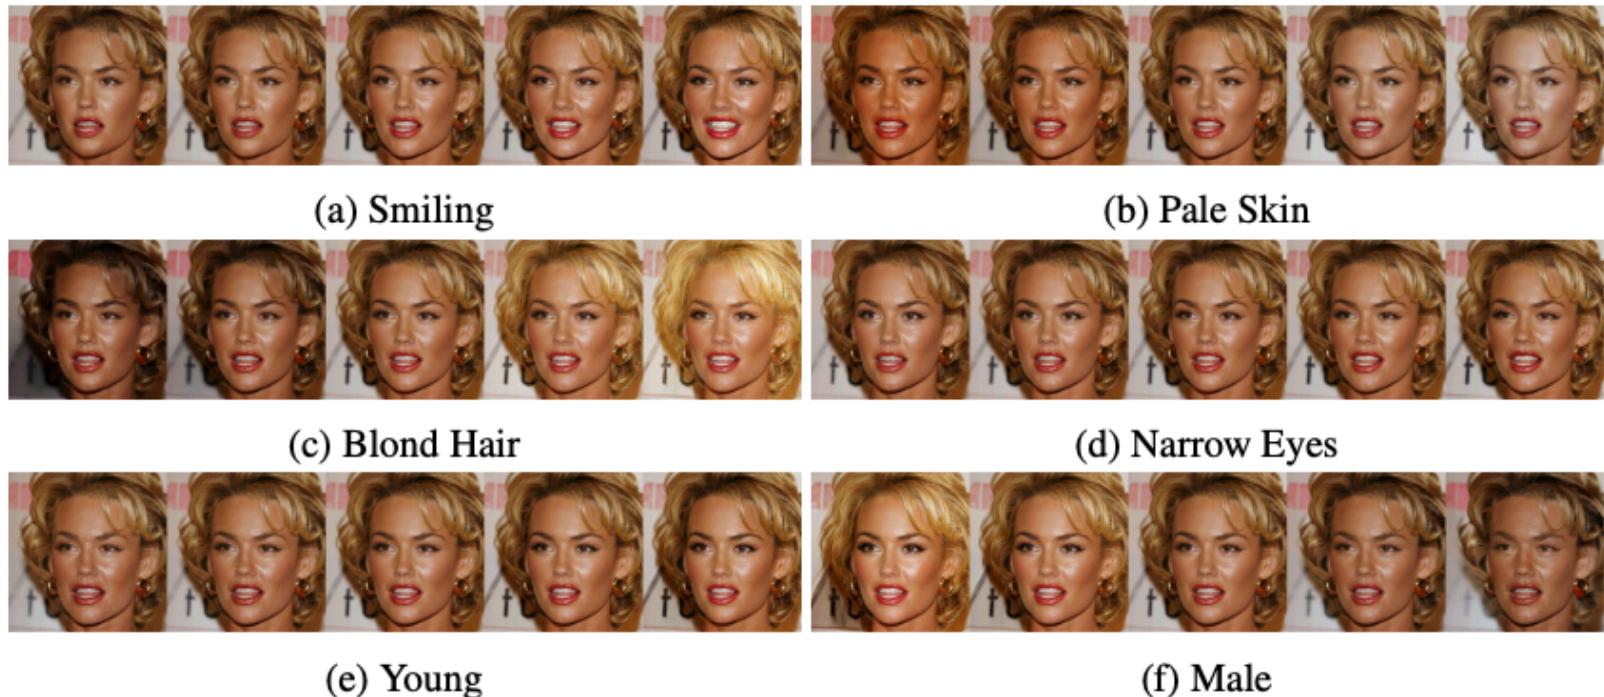


Figure 6: Manipulation of attributes of a face. Each row is made by interpolating the latent code of an image along a vector corresponding to the attribute, with the middle image being the original image

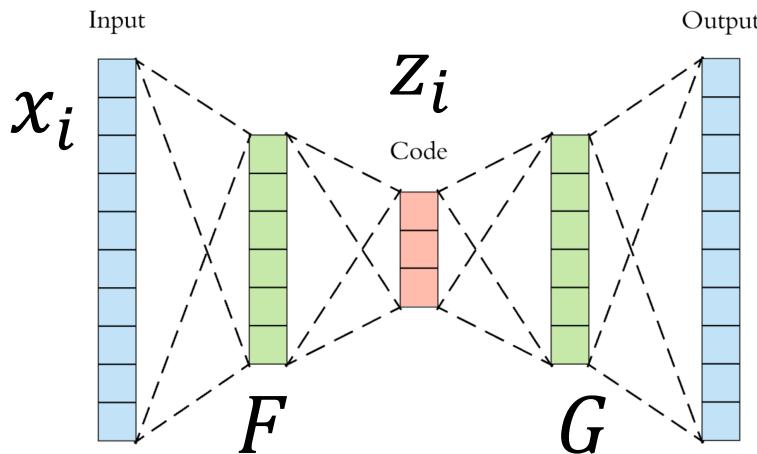
<https://papers.nips.cc/paper/8224-glow-generative-flow-with-invertible-1x1-convolutions.pdf>

Normalizing flows use invertible deep models for the generator which allow more capabilities

- ▶ Transforming between observed/input and latent space is easy
 - ▶ $x = G(z)$
 - ▶ $z = G^{-1}(x)$
- ▶ Simple sampling like GANs
 - ▶ $z \sim \text{SimpleDistribution}$
 - ▶ $x = G(z) \sim \hat{p}_g(x)$, which is estimated distribution
- ▶ **Exact density** is computable via change of variables
 - ▶ Standard maximum likelihood estimation can be used for training

Comparing VAEs and normalizing flows

VAE

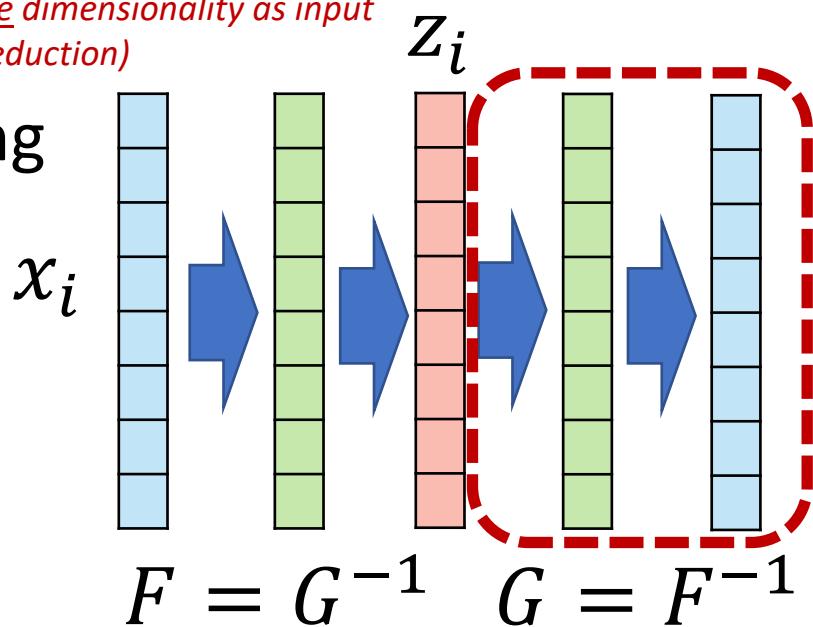


$$\tilde{x}_i \sim p(x_i | G(z_i))$$

$$L(x_i, \tilde{x}_i)$$

Latent code has same dimensionality as input
(no dimensionality reduction)

Normalizing
Flow



Implicit generator via $G = F^{-1}$
(only need to train encoder F)

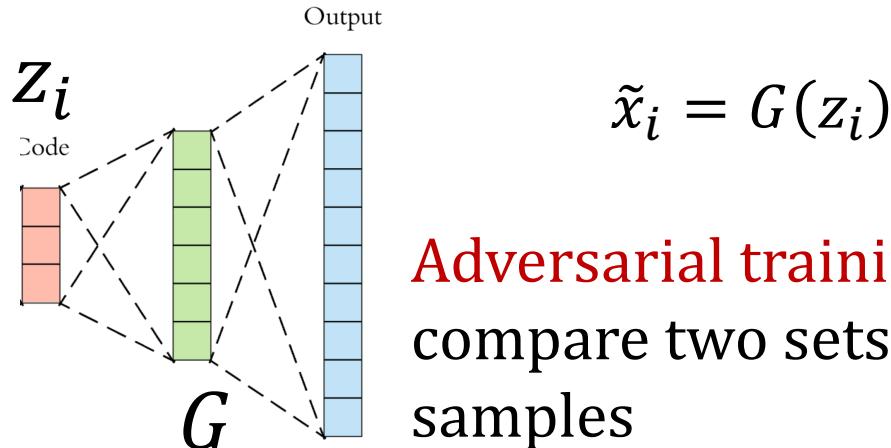
$$\begin{aligned}\tilde{x}_i &= G(F(x_i)) \\ &= G(G^{-1}(x_i)) = x_i\end{aligned}$$

$$\Rightarrow L(x_i, \tilde{x}_i) = L(x_i, x_i) = 0$$

Comparing GANs and normalizing flows:

Normalizing flows can use MLE training

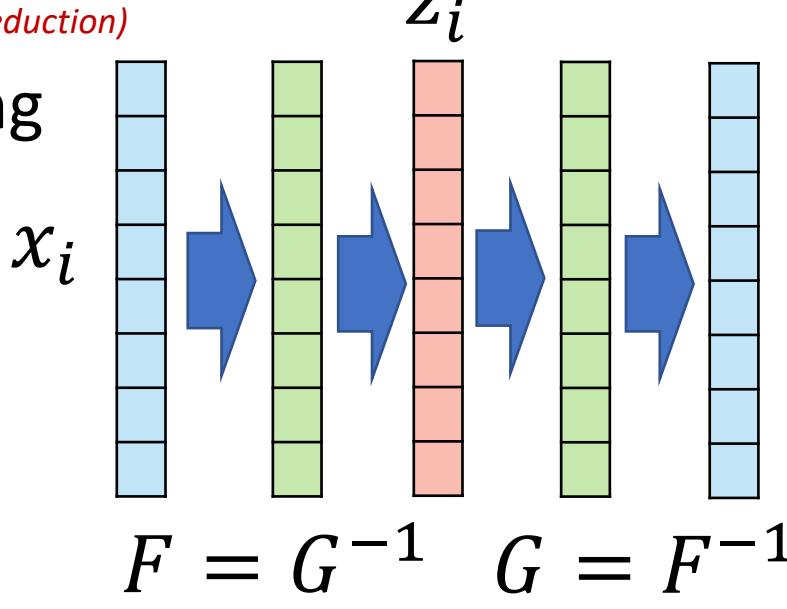
GAN



Adversarial training to compare two sets of samples

Latent code has same dimensionality as input
(no dimensionality reduction)

Normalizing Flow



$$\begin{aligned}\tilde{x}_i &= G(z_i) \\ &= G(G^{-1}(x_i)) = x_i\end{aligned}$$

MLE training since density function known

Back to maximum likelihood estimation (MLE):
How can we compute the likelihood
for normalizing flows?

- ▶ Suppose
 - ▶ $z \sim \text{Uniform}([0,1])$, i. e., $p_z(z) = 1$
(latent space is uniform)
 - ▶ $G(z) = 2z$
 - ▶ Thus, $x = G(z) = 2z$.
- ▶ What is the density function of x , what is $p_x(x)$?

Change of variables formula gives p_x in terms of the p_z and the derivative of G^{-1}

- ▶ Key idea: Must conserve density **volume** (so that distribution sums to 1).
- ▶ $p_x(x)|dx| = p_z(z)|dz|$, this is like the preservation of volume/area/mass.
 - ▶ Intuition: We only have 1 unit of “dirt” to move around.
- ▶ Rearrange above equation to get formula

$$p_x(x) = \left| \frac{dz}{dx} \right| p_z(z) = \left| \frac{dG^{-1}(x)}{dx} \right| p_z(G^{-1}(x))$$

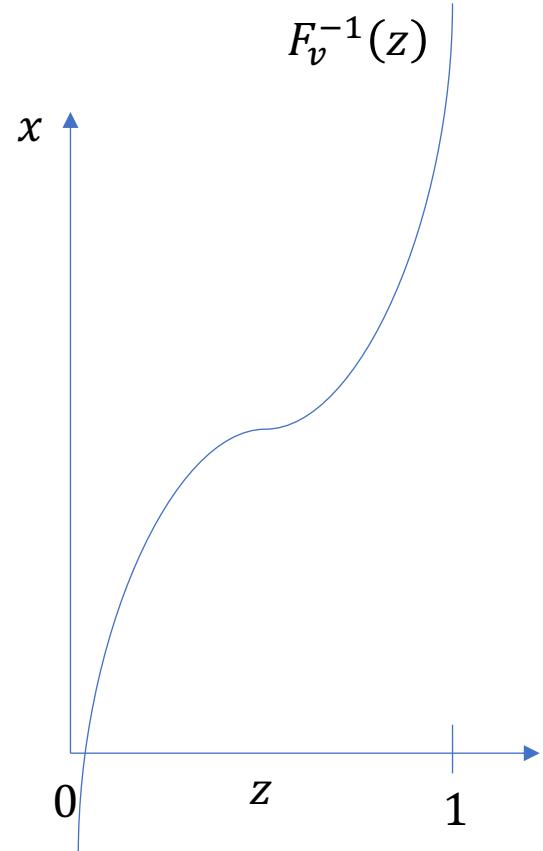
Demo of change of variables

Derivation of change of variables using CDF function (Increasing)

- ▶ Assume $x = G(z)$, where $G(z)$ is an increasing function
- ▶ $F_x(a) = \Pr(x \leq a) = \int_{-\infty}^a p_x(t)dt$
- ▶ $F_x(a) = \Pr(x \leq a)$
- ▶ $= \Pr(G(z) \leq a)$
- ▶ $= \Pr(z \leq G^{-1}(a)) = F_z(G^{-1}(a))$
- ▶ Now take the derivative of both sides with respect to a
- ▶ Left hand side: $\frac{dF_x(a)}{da} = p_x(a)$
- ▶ Right hand side: $\frac{dF_z(G^{-1}(a))}{da} = \frac{dF_z(G^{-1}(a))}{d(G^{-1}(a))} \left(\frac{dG^{-1}(a)}{da} \right)$
- ▶ $= p_z(G^{-1}(a)) \left(\frac{dG^{-1}(a)}{da} \right)$
- ▶ $\Rightarrow p_x(a) = p_z(G^{-1}(a)) \left(\frac{dG^{-1}(a)}{da} \right)$
- ▶ You can do similarly for decreasing functions to get: $p_x(a) = p_z(G^{-1}(a)) \left| \frac{dG^{-1}(a)}{da} \right|$

Inverse transform sampling
is based on change of variables

- ▶ $z \sim \text{Uniform}([0,1])$
- ▶ $v \sim \text{AnotherDistribution}$
- ▶ $x = F_v^{-1}(z)$, where F_v^{-1} is the inverse CDF for v
- ▶ What is the distribution of x ?
- ▶ $p_x(x) = p_z(F_v(x)) \left| \frac{dF_v(x)}{dx} \right|$
- ▶ $p_x(x) = (1)|p_v(x)| = p_v(x)$



What about change of variables in higher dimensions?

- ▶ Let's again build a little intuition (see demo)
- ▶ Again, conservation of volume: Consider infinitesimal expansion or shrinkage of volume
$$p(x_1, x_2) |dx_1 dx_2| = p(z_1, z_2) |dz_1 dz_2|$$
- ▶ Given that Jacobian is all mixed derivatives we get generalization for vector to vector invertible functions:

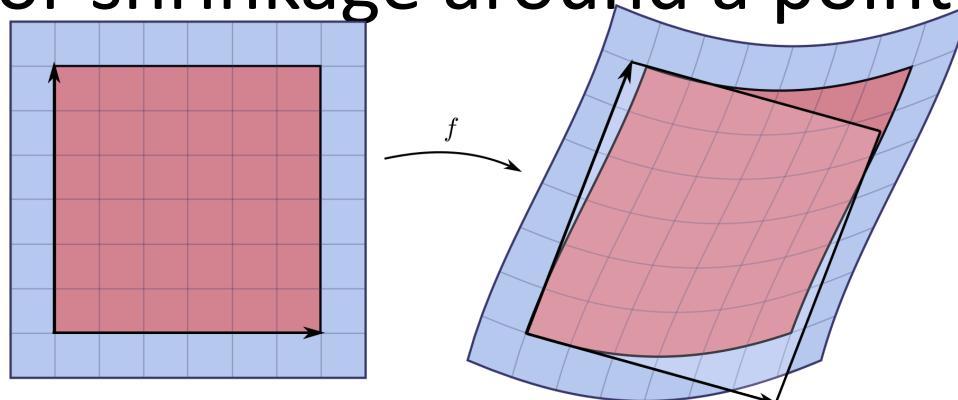
$$p_x(x) = |\det J_{G^{-1}}(x)| p_z(G^{-1}(x))$$

What is the Jacobian again?
The best linear approximation at a point

- The Jacobian definition:

$$\frac{dz}{dx} = J_z(x) = \begin{bmatrix} \frac{\partial z_1}{\partial x_1} & \dots & \frac{\partial z_1}{\partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial z_d}{\partial x_1} & \dots & \frac{\partial z_d}{\partial x_d} \end{bmatrix} = \begin{bmatrix} \frac{\partial G^{-1}(x)_1}{\partial x_1} & \dots & \frac{\partial G^{-1}(x)_1}{\partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial G^{-1}(x)_d}{\partial x_1} & \dots & \frac{\partial G^{-1}(x)_d}{\partial x_d} \end{bmatrix}$$

- The determinant measures the local *linear* expansion or shrinkage around a point



The determinant Jacobian of compositions of functions is the product of determinant Jacobians

- ▶ Suppose $F(x) = F_2(F_1(x))$
- ▶ The Jacobian expands like the chain rule

$$J_F(x) = J_{F_2}(F_1(x))J_{F_1}(x) = J_{F_2}J_{F_1}$$

- ▶ If we take the determinant of the Jacobian, then it becomes a product of determinants

$$\det J_F = \det J_{F_2}J_{F_1} = (\det J_{F_2})(\det J_{F_1})$$

- ▶ This will be useful since each layer of our flows will be invertible

Okay, now back to learning flows:
The log likelihood is the sum of determinant terms for each layer

- ▶ Simply optimize the minimize negative log likelihood where $F_\theta = G^{-1}$

$$\arg \min_{F_\theta} -\frac{1}{n} \sum_i \log p_x(x_i; \theta)$$

- ▶ $-\frac{1}{n} \sum_i \log p_z(F_\theta(x_i)) |\det J_{F_\theta}(x_i)|$
- ▶ $-\frac{1}{n} \sum_i [\log p_z(F_\theta(x_i)) + \log |\det J_{F_\theta}(x_i)|]$
- ▶ $-\frac{1}{n} \sum_i [\log p_z(F_\theta(x_i)) + \sum_\ell \log |\det J_{F_\theta^{(\ell)}}(z_i^{(\ell-1)})|]$

where $z_i^0 = x_i$, and $z_i^\ell = F_\theta^{(\ell)}(z_i^{\ell-1})$

How do we create these invertible layers?

- ▶ Consider arbitrary invertible transformation F_θ
 - ▶ How often would $|\det J_{F_\theta}|$ need to be computed?
- ▶ High computation costs
 - ▶ Determinant costs roughly $O(d^3)$ even if Jacobian is already computed!
 - ▶ Would need to be computed every stochastic gradient iteration

How do we create these invertible layers?
Independent transformation on each dimension

- ▶ $z_1 = F_1(x_1)$
- ▶ $z_2 = F_2(x_2)$
- ▶ $z_3 = F_3(x_3)$
- ▶ What is the Jacobian?

$$J_F = \begin{bmatrix} \frac{dF_1(x_1)}{dx_1} & 0 & 0 \\ 0 & \frac{dF_2(x_2)}{dx_2} & 0 \\ 0 & 0 & \frac{dF_3(x_3)}{dx_3} \end{bmatrix}$$

How do we create these invertible layers? Autoregressive Flows based on chain rule

- ▶ Forward - Density estimation (in parallel)
 - ▶ $z_1 = F_1(x_1)$
 - ▶ $z_2 = F_2(x_2|x_1)$
 - ▶ $z_3 = F_3(x_3|x_1, x_2)$
- ▶ Inverse – Sampling (conditioned on x so must be **sequential**)
 - ▶ $x_1 = F_1^{-1}(z_1)$
 - ▶ $x_2 = F_2^{-1}(z_2|x_1)$
 - ▶ $x_3 = F_3^{-1}(z_3|x_1, x_2)$
- ▶ What is the Jacobian and determinant?
 - ▶ Product of diagonal!

$$J_F = \begin{bmatrix} \frac{dF_1}{dx_1} & 0 & 0 \\ \frac{dF_2}{dx_1} & \frac{dF_2}{dx_2} & 0 \\ \frac{dF_3}{dx_1} & \frac{dF_3}{dx_2} & \frac{dF_3}{dx_3} \end{bmatrix}$$

Rezende, D., & Mohamed, S. (2015, June). Variational Inference with Normalizing Flows. In *International Conference on Machine Learning* (pp. 1530-1538).

How do we create these invertible layers? Inverse Autoregressive Flows based on chain rule

- ▶ Forward - Density estimation (**sequential**)

- ▶ $z_1 = F_1(x_1)$
- ▶ $z_2 = F_2(x_2 | z_1)$
- ▶ $z_3 = F_3(x_3 | z_1, z_2)$

- ▶ Inverse – Sampling (**parallel**)

- ▶ $x_1 = F_1^{-1}(z_1)$
- ▶ $x_2 = F_2^{-1}(z_2 | z_1)$
- ▶ $x_3 = F_3^{-1}(z_3 | z_1, z_2)$

- ▶ What is the Jacobian and determinant?
 - ▶ Product of diagonal!

$$J_F = \begin{bmatrix} \frac{dF_1}{dx_1} & 0 & 0 \\ \frac{dF_2}{dx_1} & \frac{dF_2}{dx_2} & 0 \\ \frac{dF_3}{dx_1} & \frac{dF_3}{dx_2} & \frac{dF_3}{dx_3} \end{bmatrix}$$

Kingma, D. P., Salimans, T., Jozefowicz, R., Chen, X., Sutskever, I., & Welling, M. (2016). Improved variational inference with inverse autoregressive flow. In *Advances in neural information processing systems* (pp. 4743-4751).

Scale-and-shift simple form of invertible functions (MAF <https://arxiv.org/pdf/1705.07057.pdf>)

- ▶ Forward – Density estimation (**parallel**)
 - ▶ $z_1 = \exp(\alpha_1)x_1 + \mu_1$
 - ▶ $z_2 = \exp(\alpha_2)x_2 + \mu_2, \alpha_2 = f_2(x_1), \mu_2 = g_2(x_1)$
 - ▶ $z_3 = \exp(\alpha_3)x_3 + \mu_3, \alpha_3 = f_3(x_1, x_2), \mu_3 = g_3(x_1, x_2)$
- ▶ What is the Jacobian and determinant?

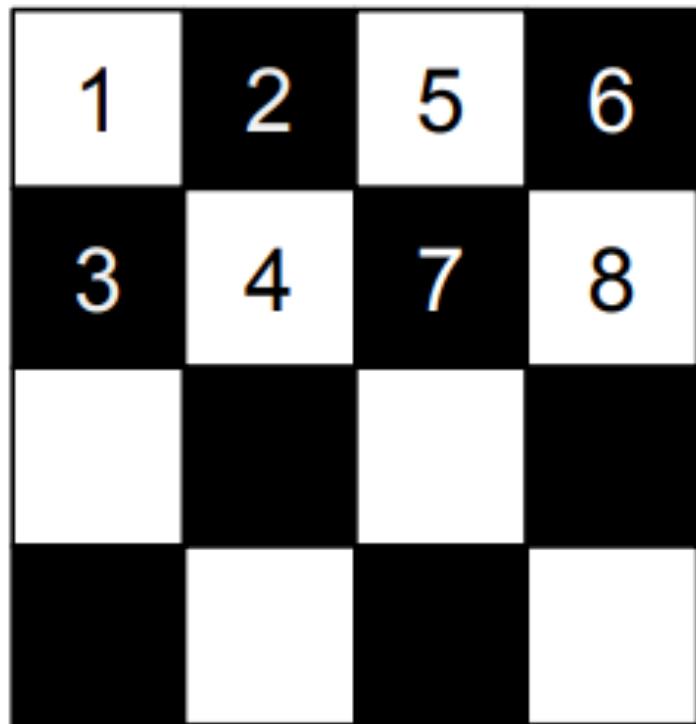
$$J_F = \begin{bmatrix} \exp(\alpha_1) & 0 & 0 \\ \frac{dz_2}{dx_1} & \exp(\alpha_2) & 0 \\ \frac{dz_3}{dx_1} & \frac{dz_3}{dx_2} & \exp(\alpha_3) \end{bmatrix}$$

Coupling layers allow parallel density estimation and sampling

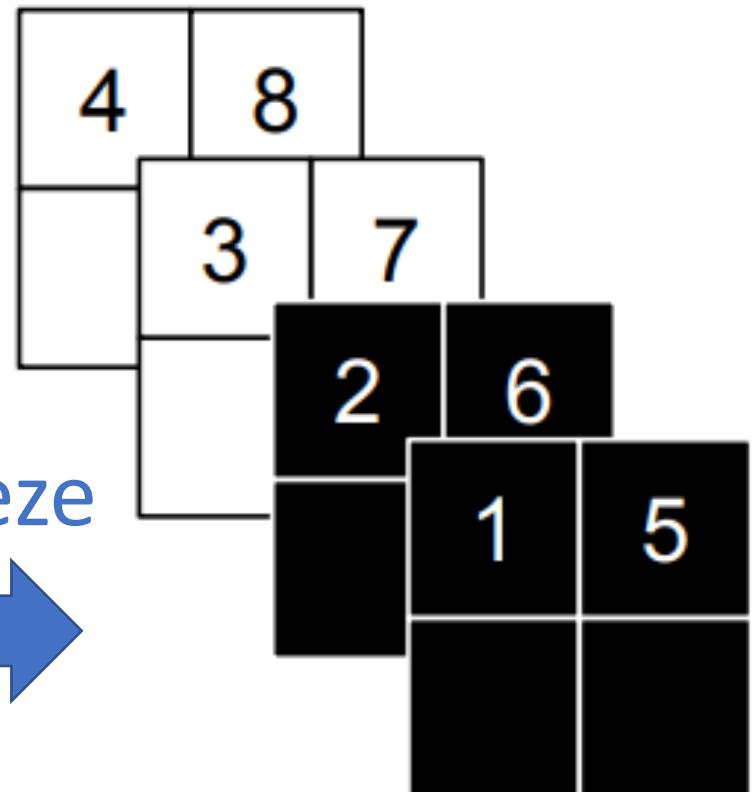
- ▶ Keep some set of features **fixed** and transform others
 - ▶ $z_{1:i-1} = x_{1:i-1}$
 - ▶ $z_{i:d} = \exp(f(x_{1:i-1})) \odot x_{i:d} + g(x_{1:i-1})$
- ▶ Reverse or shuffle coordinates and repeat
- ▶ What is Jacobian?

$$J_F = \begin{bmatrix} I & 0 \\ J_{cross} & \text{diag}(\exp(f(x_{1:i-1}))) \end{bmatrix}$$

The squeeze operation trades off between spatial and channel dimensions



Squeeze



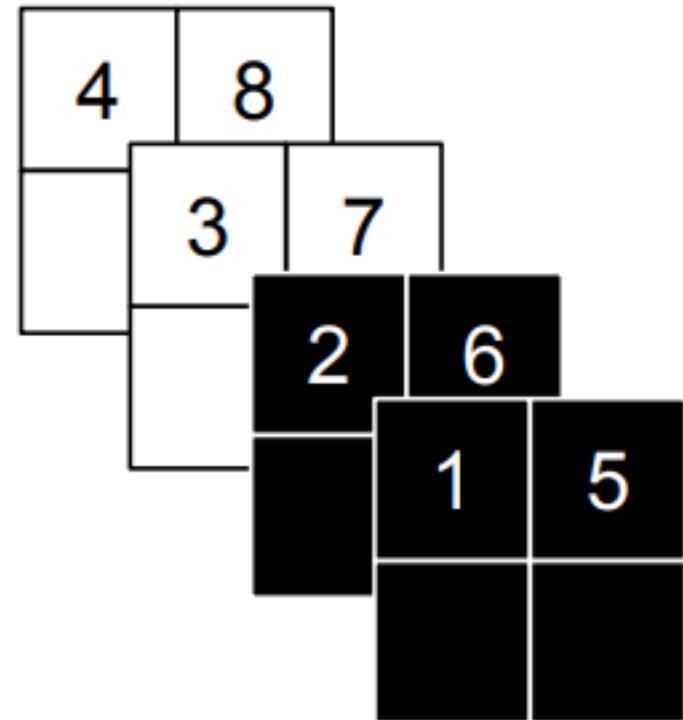
$H \times W \times C$

$H/2 \times W/2 \times 4C$

Dinh, L., Sohl-Dickstein, J., & Bengio, S. (2016). Density estimation using real nvp. *arXiv preprint arXiv:1605.08803*.

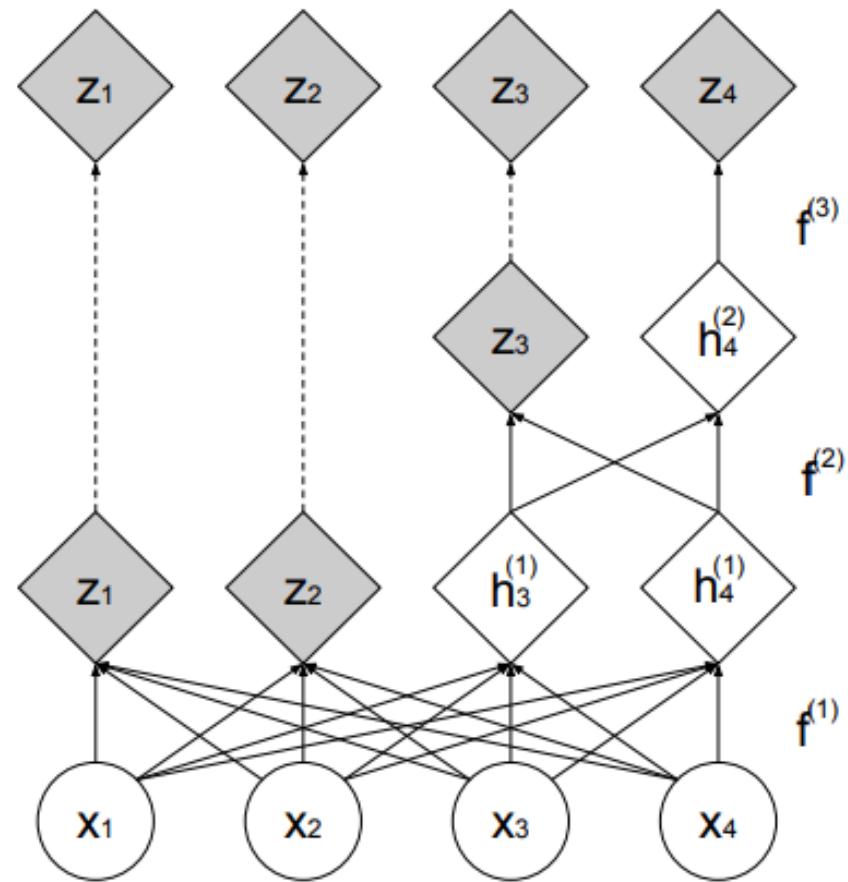
Checkboard or channel-wise masking can be used to separate fixed and non-fixed set of variables

White are **fixed**, i.e., $x_{1:i-1}$, and black are **transformed**, $x_{i:d}$.



Hierarchical factorization is like an invertible dimensionality reduction method

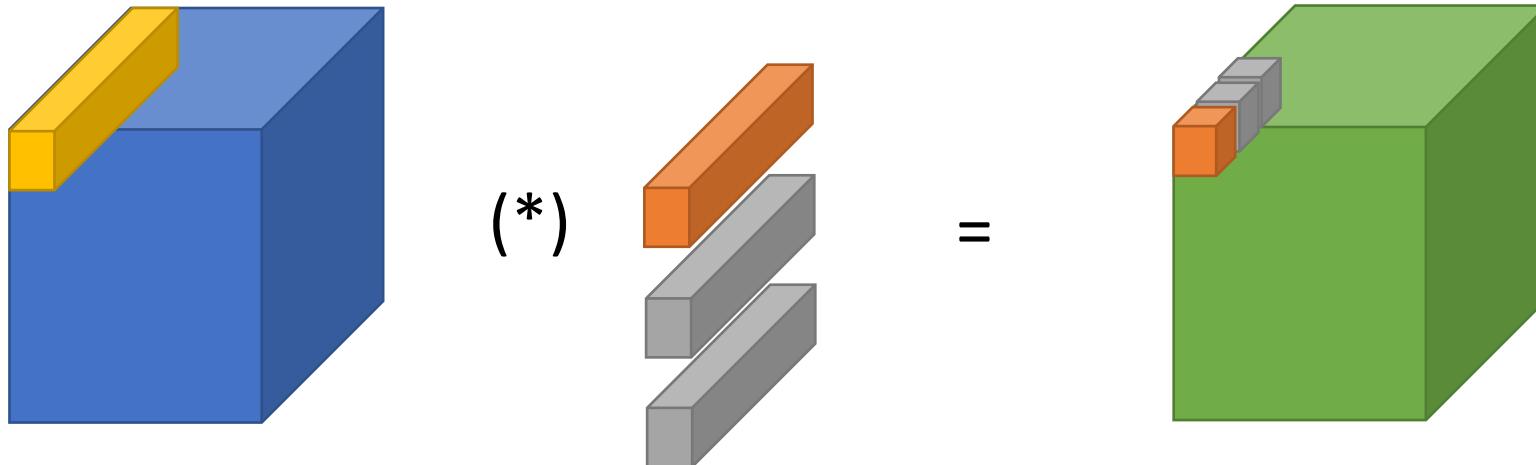
- ▶ After each block, half of the dimensions are fixed and the rest pass through more transformations
- ▶ Intuitively, the important part of the signal propagates deeper



GLOW: Convolutional flows

1×1 invertible convolutions are like fully connected layers for each pixel

- ▶ Image tensor: $h \times w \times c$
- ▶ If we use c filters than we map from a $h \times w \times c$ to another $h \times w \times c$ image
- ▶ The number of parameters is a matrix $W \in \mathbb{R}^{c \times c}$
- ▶ 1×1 convolutions can be seen as a linear transform along the channel dimension (mixes the channel dimensions)



Highly realistic random samples from powerful flow model (GLOW)



Figure 1: Synthetic celebrities sampled from our model; see Section 3 for architecture and method, and Section 5 for more results.

<https://papers.nips.cc/paper/8224-glow-generative-flow-with-invertible-1x1-convolutions.pdf>

Similar concepts can be used to generate realistic audio (WaveGlow)

- ▶ Listen to some examples

<https://nv-adlr.github.io/WaveGlow>

- ▶ Very similar concepts for audio generation