

Spring 2022

ECE568:

Software for Embedded

Systems

Lecture #5: Embedded Software Architectures

Vijay Raghunathan
School of ECE, Purdue University
vr@purdue.edu

(c) 2022, All rights reserved.

Outline

2

- Single Program Approach (SPA)
- Foreground-Background Systems
- Multi-tasking Systems
 - Event driven architectures
 - Thread based architectures
- An anatomy of two simple embedded operating systems
- How can we organize / schedule computation in these various architectures?

A Simple Embedded System

3

- As an example, let's consider a simple embedded system with a keypad, an LCD, and an RS-232 port that runs some communications. The system also has some GPIO and a parallel printer. Each change of state of an input or output results in an RS-232 message sent out and an LCD update. Received RS-232 messages can result in printouts, LCD updates, and output status updates. We may have to start a flash pattern on a particular LED as a result of:
 - An input or output becoming active
 - Keypad entry
 - Communications message received

Single Program Approach:

4

```
int main(void)    {
    Init_All();
    for (;;) {
        IO_Inputs_Scan();
        KBD_Scan();
        RS232_Receive();
        IO_Process_Outputs();
        LCD_Update();
        RS232_Send();
        PRN_Print();
        Maintenance_Process();
    }
    // should never get here; put error
    // handling here, just in case
    return (0);
}
```

[Melkonian00]

Observations on the SPA

5

- Essentially an infinite loop
 - Each function called in the loop represents an independent task
- Each of these tasks must return in a reasonable time, no matter what code is being executed (who ensures this?)
- We have no idea at what frequency our main loop runs
 - The frequency is not constant and can significantly change with changes in system status (e.g., displaying a large image, printing a long document)
- Mix of periodic and event-driven tasks
 - Most tasks are event driven (have dedicated input event queues)
 - e.g., IO_Process_Outputs receives events from RS232_Receive, IO_Inputs_Scan, and KBD_Scan when an output needs to be turned on
 - Others are periodic in nature
 - No trigger event, but may have different (time-varying) periods

Observations on the SPA (contd.)

6

- Need some simple means of inter-task communications
 - e.g., may want to stop scanning the inputs after a particular keypad entry and restart the scanning after another entry
 - require a call from the keypad scanner to stop the I/O scanner task
 - e.g., may also want to slow down the execution of some tasks depending on the circumstances
 - say we detect an avalanche of input state changes, and our RS-232 link can no longer cope with sending all these messages

Observations on the SPA (contd.)

7

- Usually, each task in the loop checks for the occurrence of asynchronous events using “polling”
 - i.e., must explicitly check if an event has occurred or not
 - Depending on result of the poll, task may do different things (e.g., if keyboard input received, then process it. If no keyboard input, just return)
 - Adds to timing uncertainty of the infinite loop
 - Not very resource efficient

Foreground/Background (FB) Sys.

8

- Break workload into two levels
 - Critical tasks (tight timing constraints) form the foreground while other tasks (best-effort) form the background

Background:

```
while True:  
    if (time for display update):  
        do display  
    else if (time for operator input):  
        do operator  
    else if (time for mgmnt. request):  
        do mgmnt.
```

Foreground (interrupt):

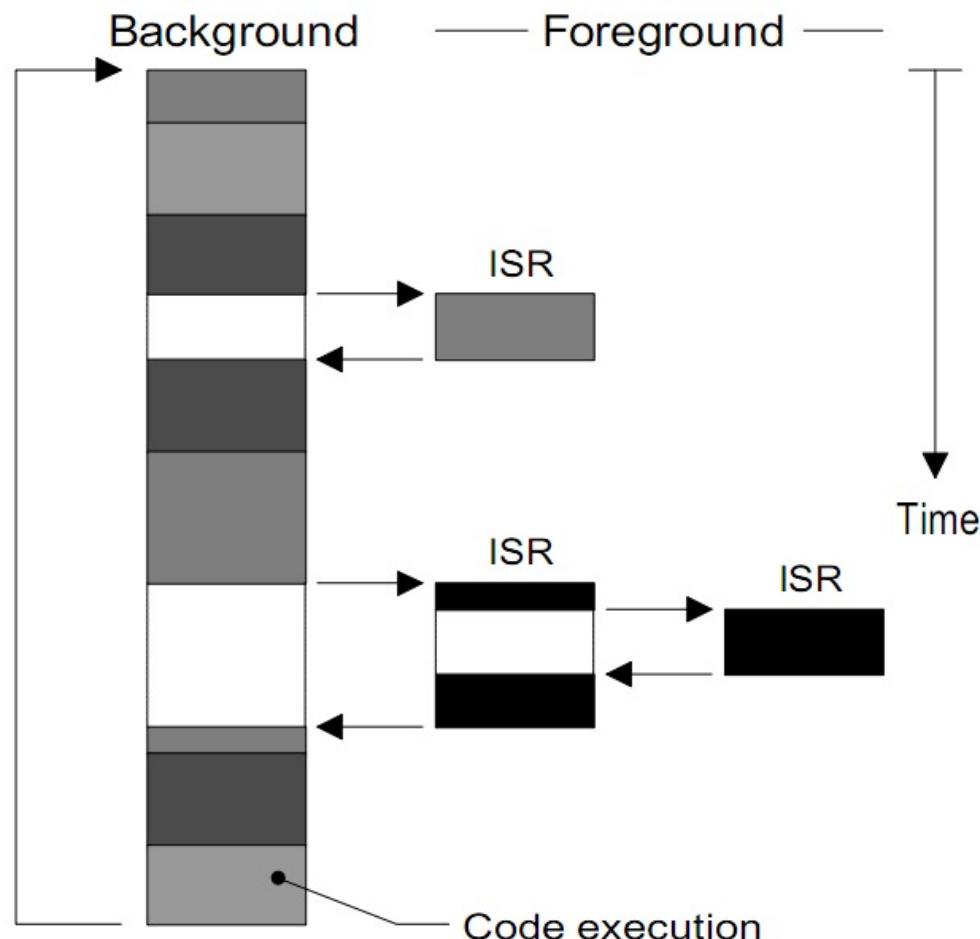
```
on interrupt  
    do clock module  
    if (time for control):  
        do control
```

- Decoupling relaxes timing constraint

Interrupts in FB Systems

9

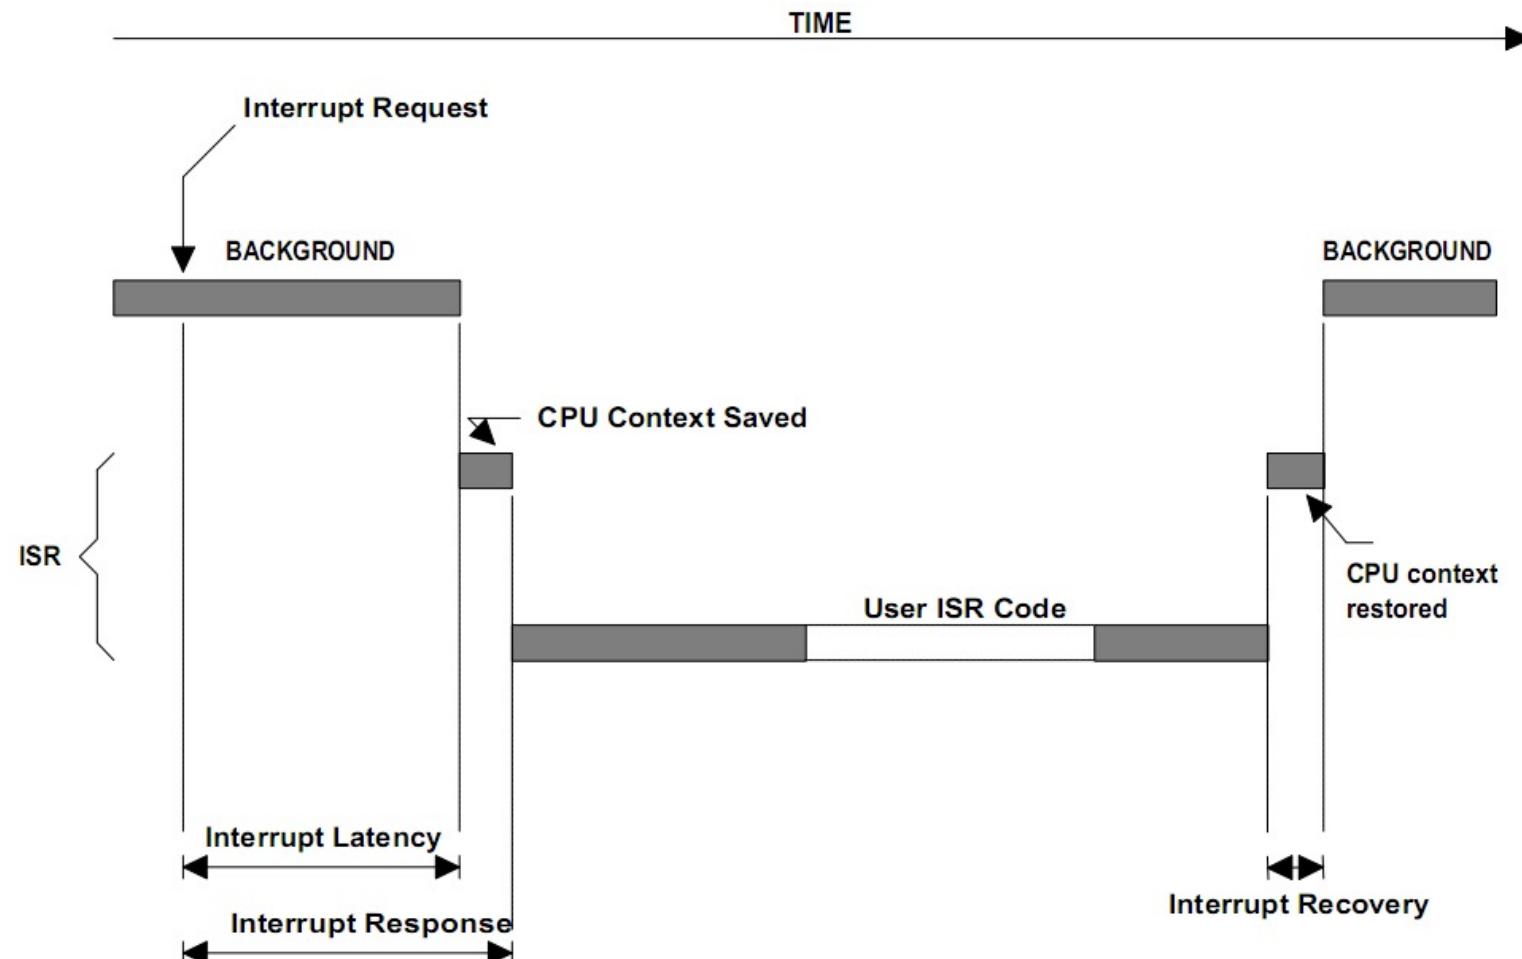
- Task-level response time depends on how fast the background loop runs and when interrupt occurs



Interrupt-related Timing Metrics

10

- Disabling interrupts affects interrupt latency
- Context switch time depends on number of registers



Multi-tasking Approach

11

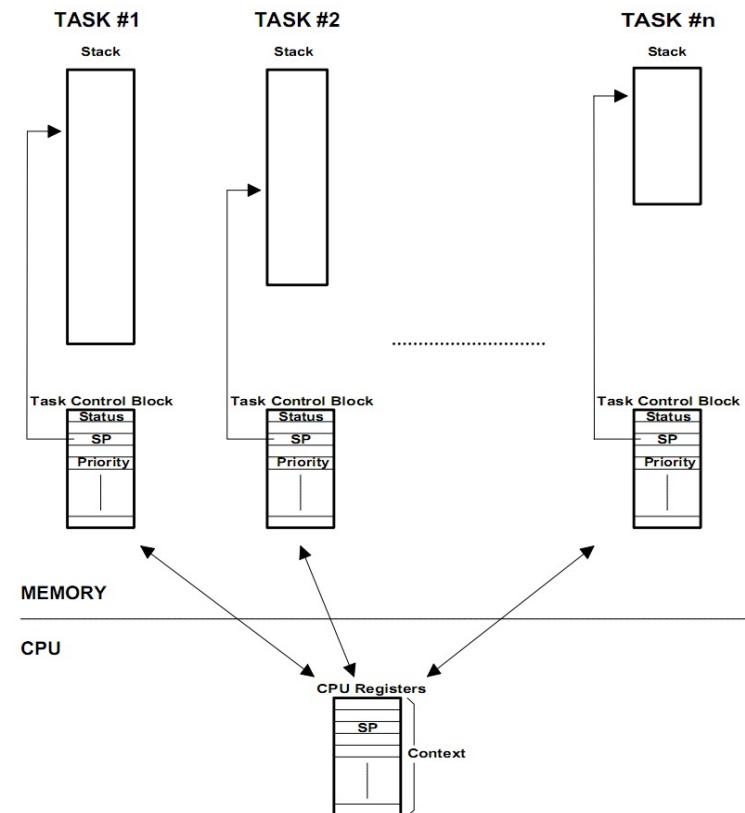
- Single program approach: one “task”
- Foreground/background: two tasks (or levels)
- Generalization: multiple tasks (processes, threads, etc.)
 - ▶ Time share the CPU between these tasks
 - ▶ Similar to FB system with multiple backgrounds
 - ▶ Each task can be carried out in parallel
 - Tasks simultaneously interact with external elements
 - Monitor sensors, control actuators, interrupts, I/O, etc.
 - Makes writing applications easier
 - ▶ Requires
 - Scheduling of these tasks (When? Which? How?)
 - Sharing data between concurrent tasks

Multi-tasking Approach

12

- **Context Switch:** Process of switching execution from one task to another
 - Context of a task represents snapshot of everything needed to restart the task later from where it was stopped
- When done fast, creates the illusion of concurrency

Context usually consists of processor registers (including program counter, stack pointer)



Task Characteristics

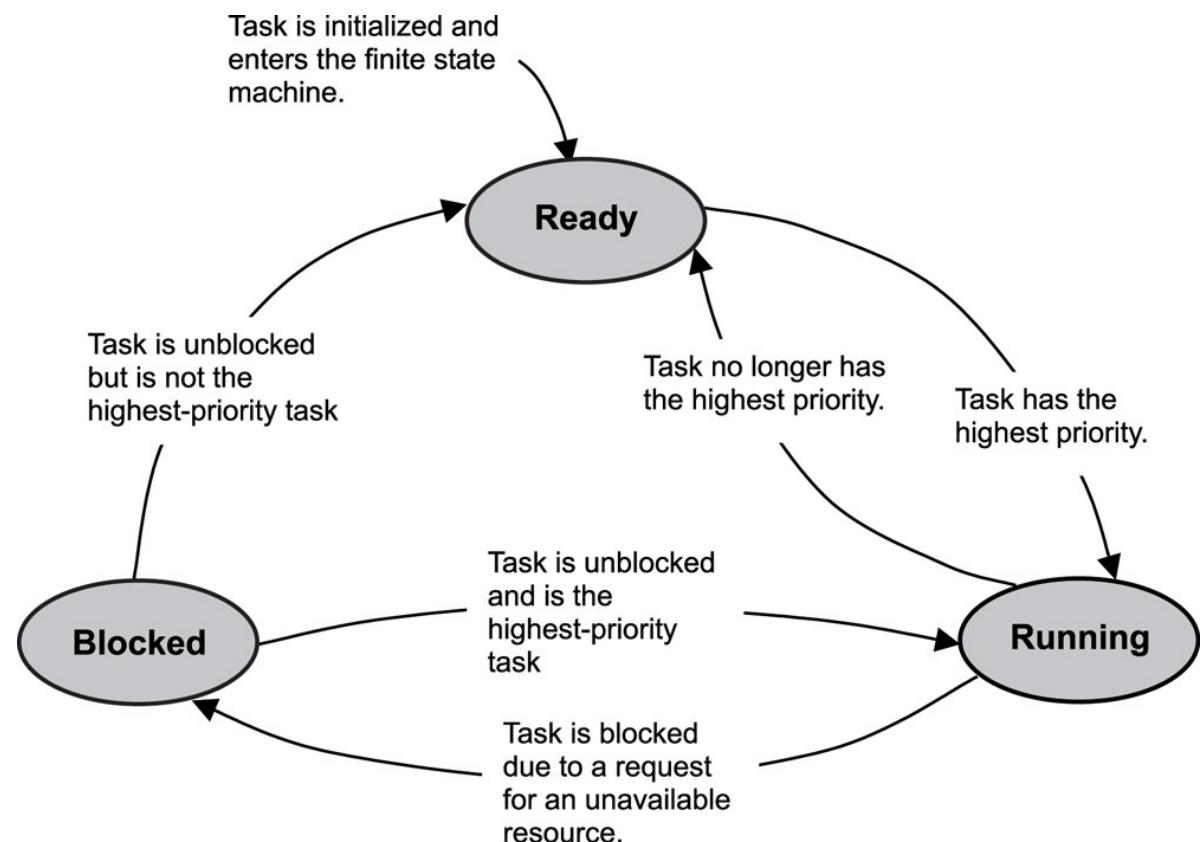
13

- New instances of tasks may be activated over time
 - Periodic vs. aperiodic (sporadic) activation
 - Also called “jobs”
- Tasks may have:
 - Resource requirements (e.g., need X clock cycles of CPU time)
 - Importance levels (e.g., priorities or criticality)
 - Precedence relationships
 - Communication requirements
 - Of course, timing constraints!
 - specify time at which task is to be performed, or is to be completed
 - e.g., period of a periodic task, deadline of a task

Task States

14

- A task can be in one of several states
 - Inactive: Not active or dormant
 - **Ready** to run: Waiting for CPU
 - Maintained in a ready queue
 - **Running**: Currently in execution
 - **Blocked**: Waiting for some event or shared resource (not CPU)
 - Interrupted: Has been interrupted to service an interrupt



Who manages the multiple tasks?

15

- Application tasks should not have to worry about
 - What happens to other tasks
 - Low-level hardware access
 - Coordinating the use of shared resources
- Have a layer of software to handle all of these and hide them from application tasks
 - Called by many names: operating system (OS), kernel, runtime, etc.
 - Provides execution environment for application tasks and various services through a well-defined interface (API)
 - Real-time OS (RTOS): Emphasis is on satisfying timing constraints (to construct real-time systems)

Separation of concerns: Let tasks focus on functionality and let the OS take care of orchestrating task execution

How to schedule multiple tasks?

16

- Cyclic executive model
 - Simple, static schedulability analysis
 - Resulting schedule or table can be used at run time
 - Time Division Multiple Access (TDMA)-like scheduling
- Event-driven model
 - Tasks are represented by functions that are **handlers for events**
 - Next event processed after function for previous event finishes
- Thread-based model
 - Preemptive vs. non-preemptive
 - Priority based scheduling: *Static* vs. *dynamic* priority
 - Often, static schedulability analysis can be done
 - No explicit schedule constructed offline. at run time tasks are executed “highest priority first”

Cyclic Executive Scheduling

20

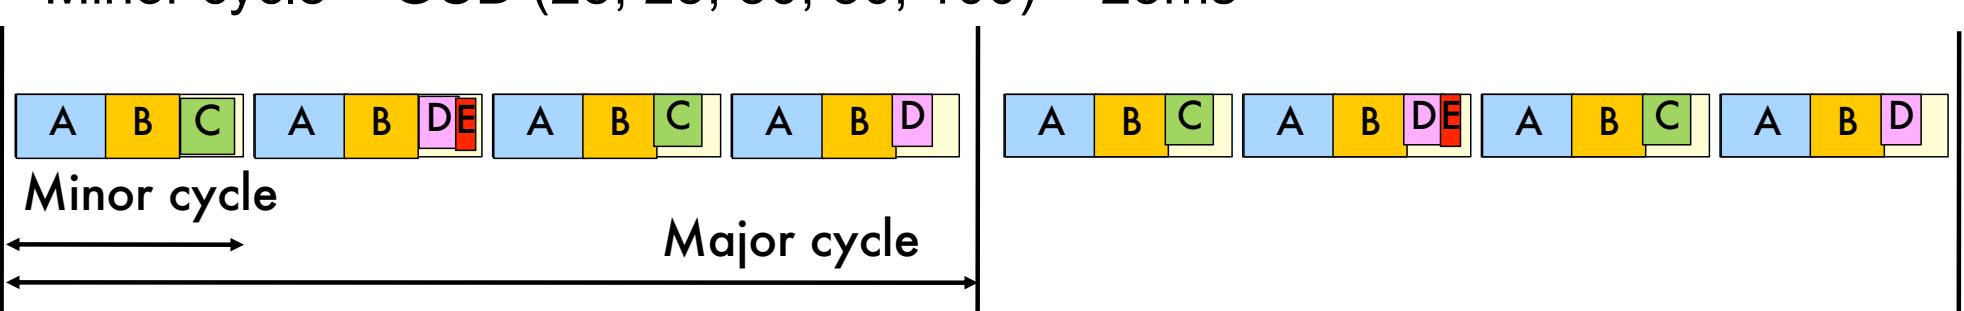
- Applicable to tasks that are periodic in nature
 - For aperiodic tasks, use worst case inter-arrival time
- Basic approach: time slices of equal length organized into major and minor cycles
 - Major cycle = LCM, Minor cycle = GCD
 - A timer with period = minor cycle interrupts a scheduler, which activates the task according to order in major cycle
 - Sum of executions within a time slice \leq minor cycle
 - May need to split tasks
 - Handling deadlines may be difficult

Cyclic Executive Example

22

Task	Time Period (P)	Exec. time (C)
A	25 ms	10 ms
B	25 ms	8 ms
C	50 ms	5 ms
D	50 ms	4 ms
E	100 ms	2 ms

- Major cycle = LCM (25, 25, 50, 50, 100) = 100ms
- Minor cycle = GCD (25, 25, 50, 50, 100) = 25ms



Cyclic Executive Example

23

```
start_up;  
loop  
    wait_for_minor_cycle_begin;  
    procedure_for_A;  
    procedure_for_B;  
    procedure_for_C;  
    wait_for_minor_cycle_begin;  
    procedure_for_A;  
    procedure_for_B;  
    procedure_for_D;  
    procedure_for_E;  
    wait_for_minor_cycle_begin;  
    procedure_for_A;  
    procedure_for_B;  
    procedure_for_C;  
    wait_for_minor_cycle_begin;  
    procedure_for_A;  
    procedure_for_B;  
    procedure_for_D;  
end loop;
```

Cyclic Executive Scheduling

26

- Advantages

- Simple to implement
- Predictable: sequence of tasks is always same
- No context switches

- Disadvantages

- Fitting code segments into slots
 - Time consuming iterative process
- Inflexible, sensitivity to application changes
 - Table is completely overhauled when tasks or their characteristics change
 - Task splitting may change
- Fragile during overload conditions
 - If a task doesn't finish on time, both letting it continue or aborting it have problems (disruption to scheduling vs. inconsistent system state)
- Modification to add functions can be expensive in redesign time > fragile design

Outline

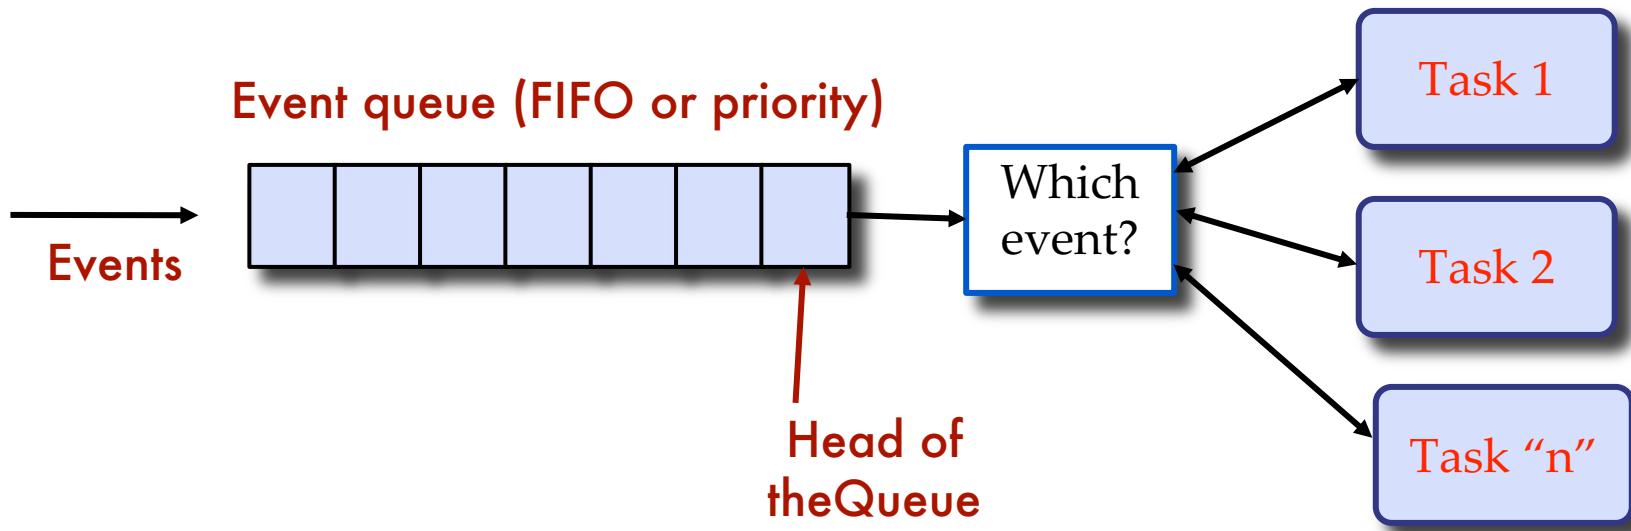
27

- Single Program Approach (SPA)
- Foreground-Background Systems
- Multi-tasking Systems
 - Event driven architectures
 - Thread based architectures
- An anatomy of two simple embedded operating systems
- How can we organize / schedule computation in these various architectures?

Event-driven Scheduling

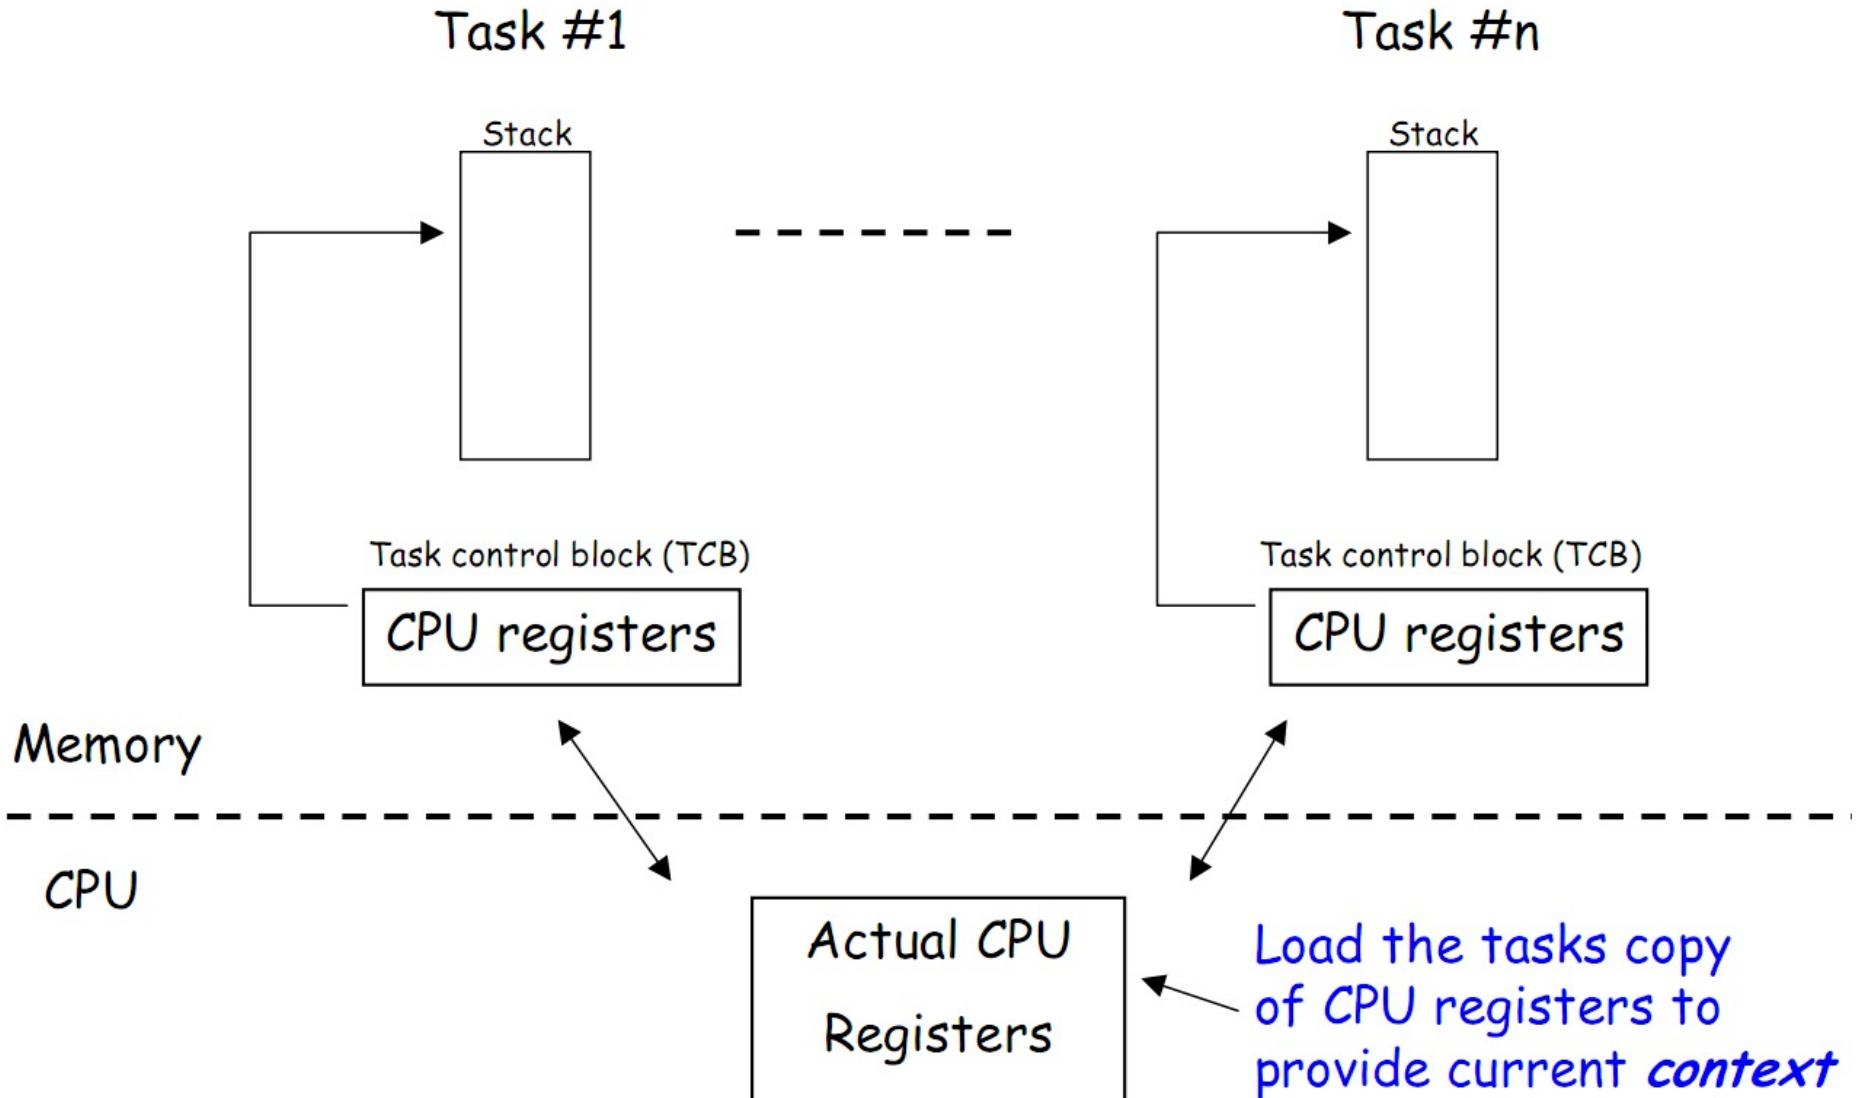
28

- System consists of tasks that respond to events
 - Events can be interrupts or software events (e.g., messages from other tasks)
- Kernel is a main loop that does the following
 - Check if event has occurred. If so, select an event to process
 - Handle the event (i.e., call task that corresponds to the event)
- Each task is an “event handler” function
 - Programmers need to create these handlers and bind them to events



Recap: Multi-tasking model

30



Preemption

31

- **Preemption:** Once a task starts, can it be stopped to execute another task before it completes?
 - Interrupts (when not disabled) can always stop tasks
- **Non-preemptive systems:** task, once started, runs until it ends or has to do some I/O
 - Lower number of context switches
- **Preemptive system:** task may be stopped to run another
 - Incurs overhead and implementation complexity
 - But has better schedulability and easier analysis

Illustration: Non-preemptive system

32

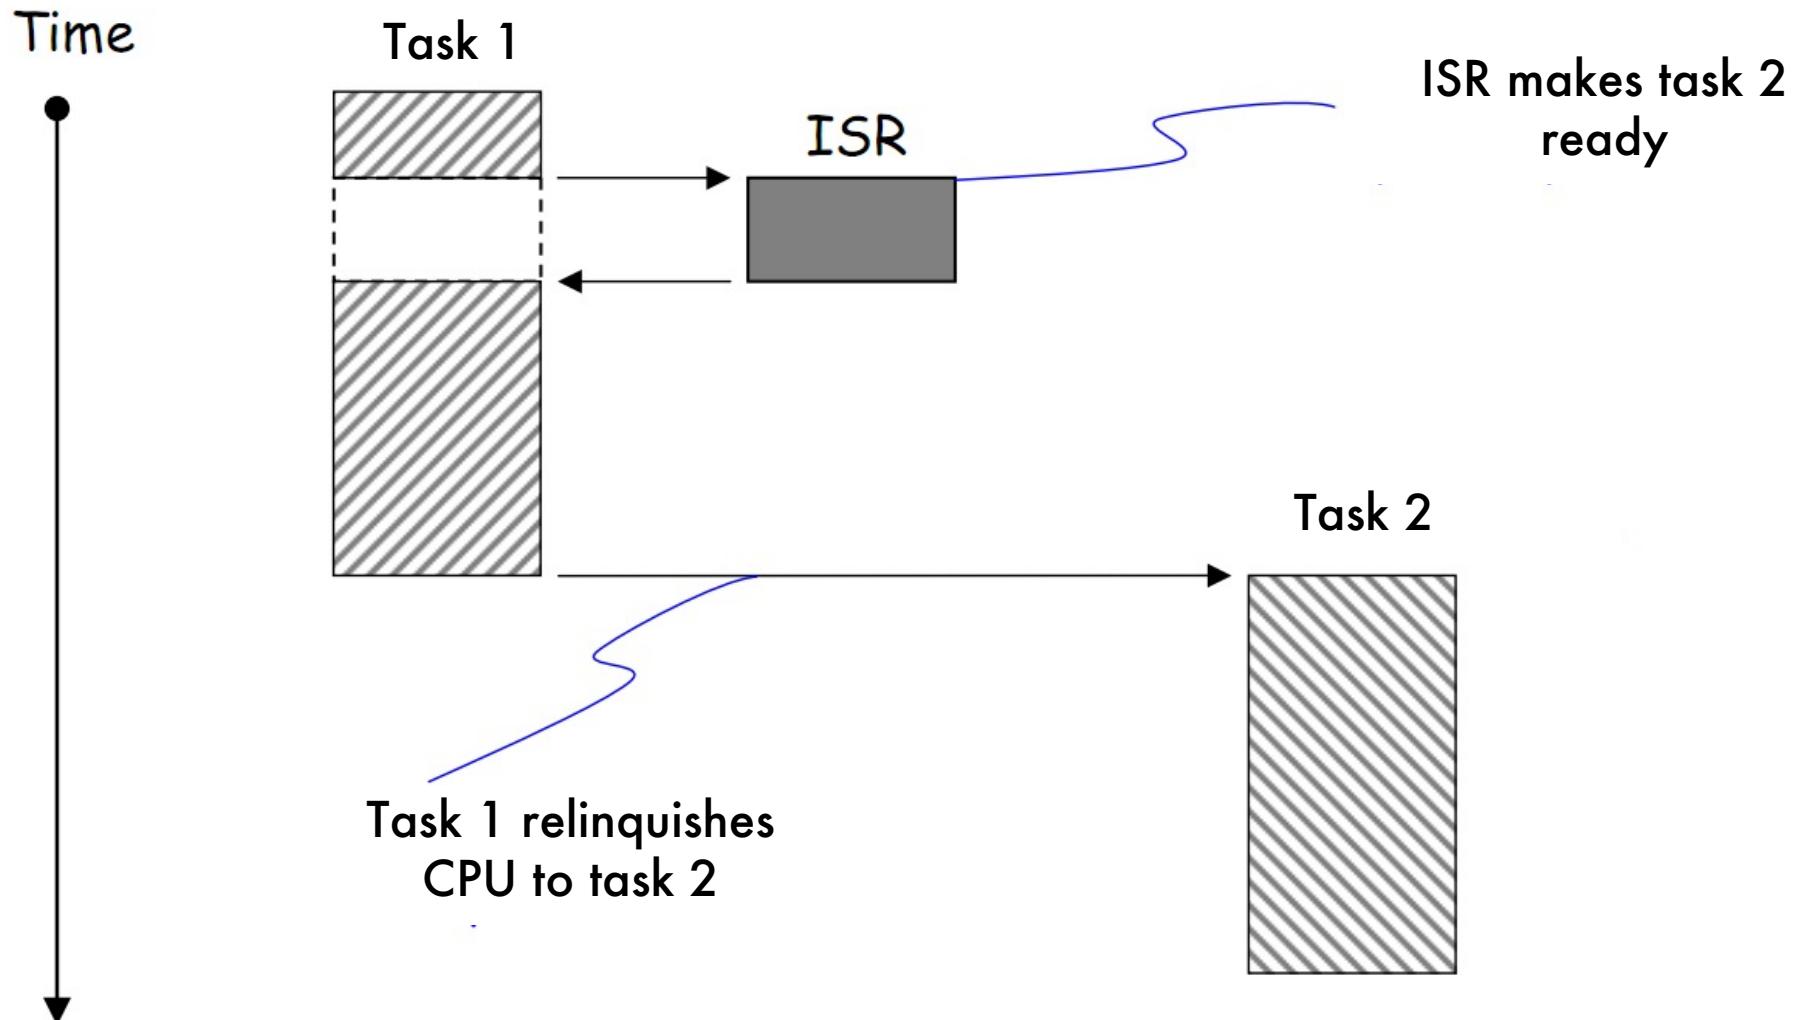
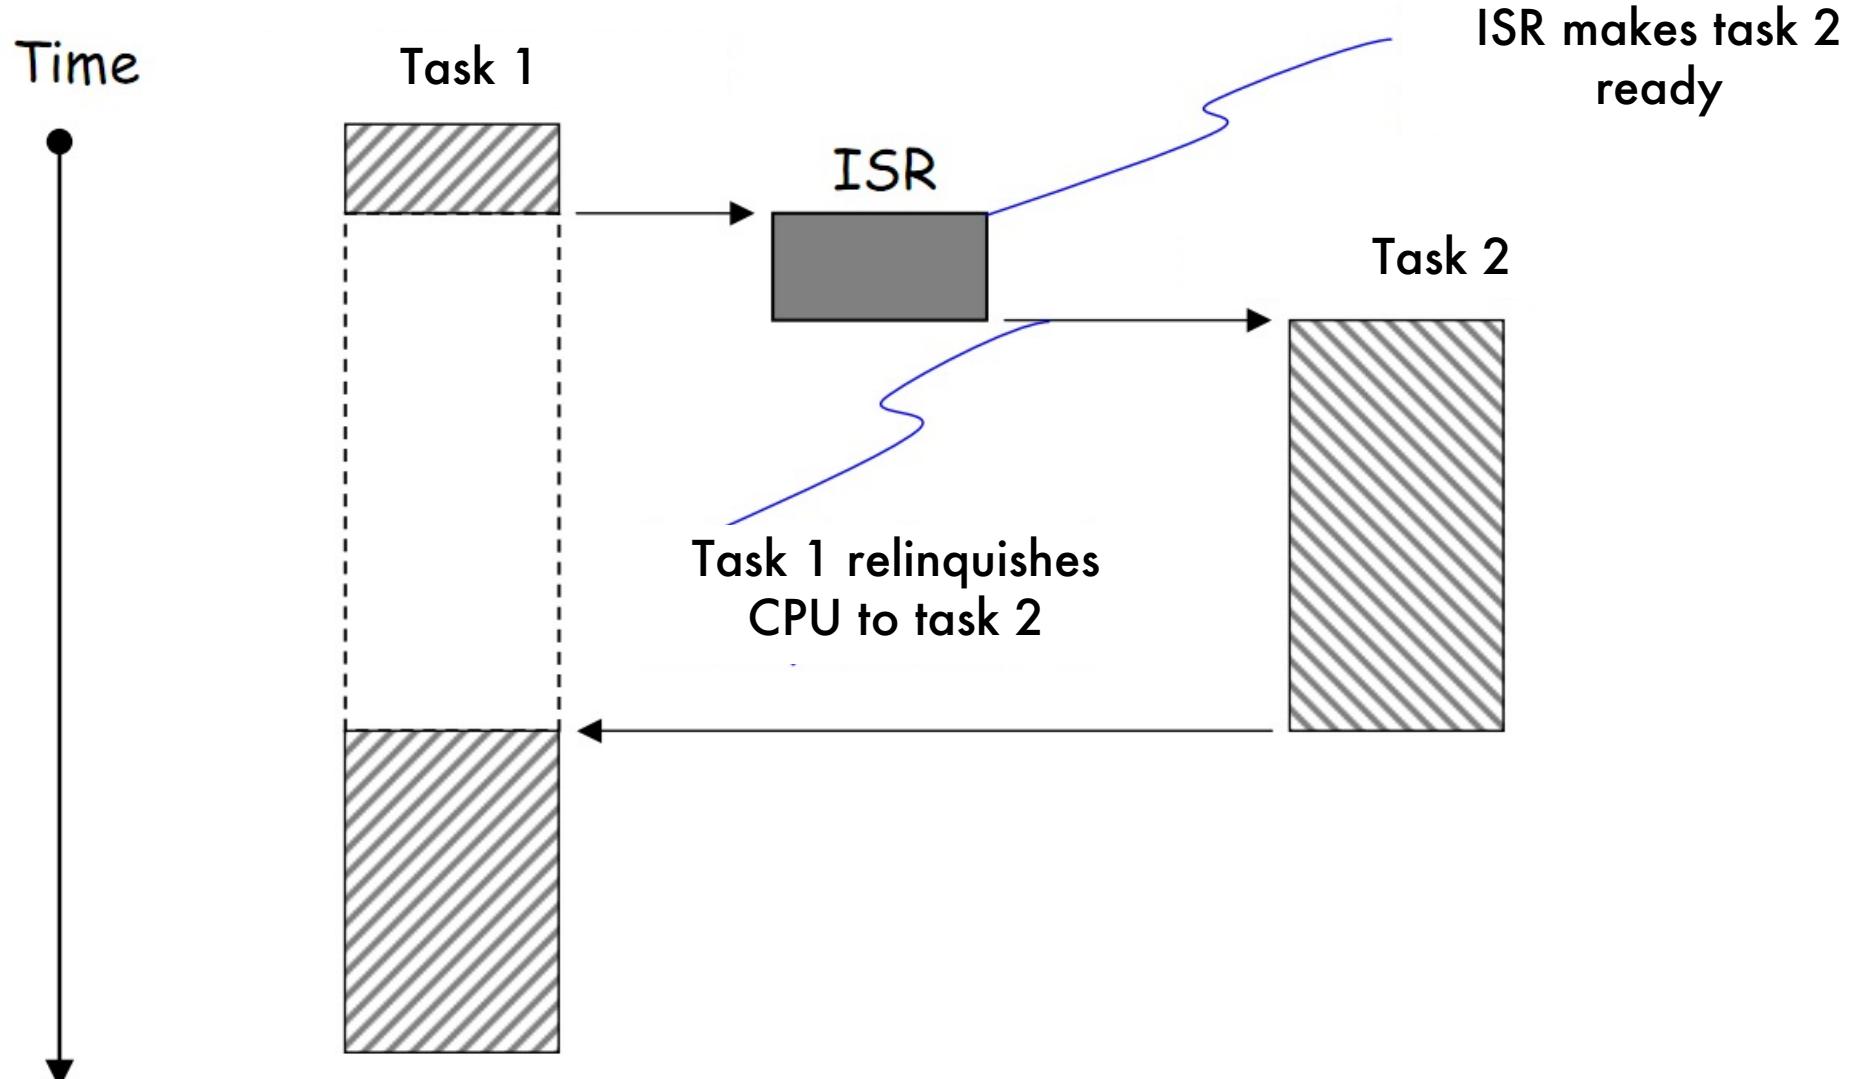


Illustration: Preemptive system

33



Priority-based Scheduling

34

- Tasks assigned priorities, statically or dynamically
 - Priority assignments relate to timing constraints
 - Static priority attractive... no recalculation, cheap
- Can be preemptive, non-preemptive, or time-sliced
- Appropriate assignment of priorities enable satisfaction of timing constraints
- Example
 - Task 1: Time Period = 2s, Requires 1s of CPU; Deadline = 2s
 - Task 2: Time Period = 5s, Requires 2s of CPU; Deadline = 5s
 - Consider: Priority1 > Priority2
 - Consider: Priority 2 > Priority 1
 - Which works?

Some Priority-based Approaches

35

- Rate Monotonic (RM) algorithm
 - Static priorities based on task periods
 - higher priorities to shorter periods
 - Optimal among all static priority schemes
- Earliest Deadline First (EDF) algorithm
 - Dynamic priority assignment
 - Closer a task's deadline, higher is its priority
 - Applicable to both periodic and aperiodic tasks
 - Optimal among all scheduling schemes for 1 CPU
 - Need to calculate priorities when new tasks arrive
 - more expensive in terms of run-time overheads
- For both, key results on schedulability exist

Embedded and Real Time OSs

38

- Usually one of four varieties
 - Small, fast, proprietary kernels
 - Real-time extensions to commercial OSs
 - Research OSs
 - Part of language run-time environments, e.g., Java (embedded, real-time Java)
- Mainly differ with respect to the following main issues
 - Concurrency model: event, thread, scheduling, etc.
 - Memory model: static vs. dynamic
 - Component model, if any

Some RTOS Examples

43

- For tiny systems
 - TinyOS, SOS, Contiki, NuttOS, ERIKA, FreeRTOS
- For mid-size systems
 - μCOS-III, eCos, OSE
 - Embedded Linuxes: μCLinux
 - Inferno
- For large-size systems
 - Windriver VxWorks
 - QNX Neutrino
 - RTEMS
 - Real-time Linuxes: RTLinux, RTAI, Linux/RK
 - Research kernels: SHARK, MARTE

Next: Two RTOS examples

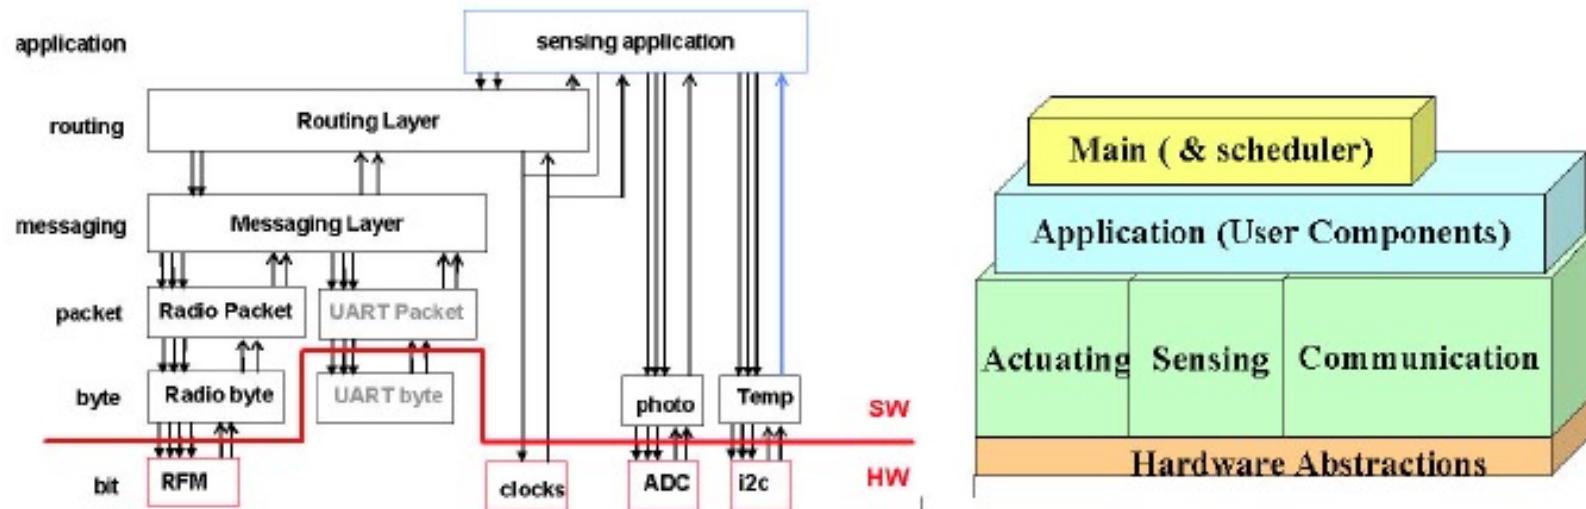
49

- A tour through the internals of two contrasting OSs
 - TinyOS
 - Event driven model
 - uC/OS-II
 - Thread based model

Example: TinyOS

50

- Dominant OS for resource-constrained wireless sensor nodes
- Not really an OS in the traditional sense
 - More like an application-specific operating system
 - Application = Scheduler + Graph of Components
- Event-driven programming model
- Static memory allocation, program analysis



Building Blocks of TinyOS

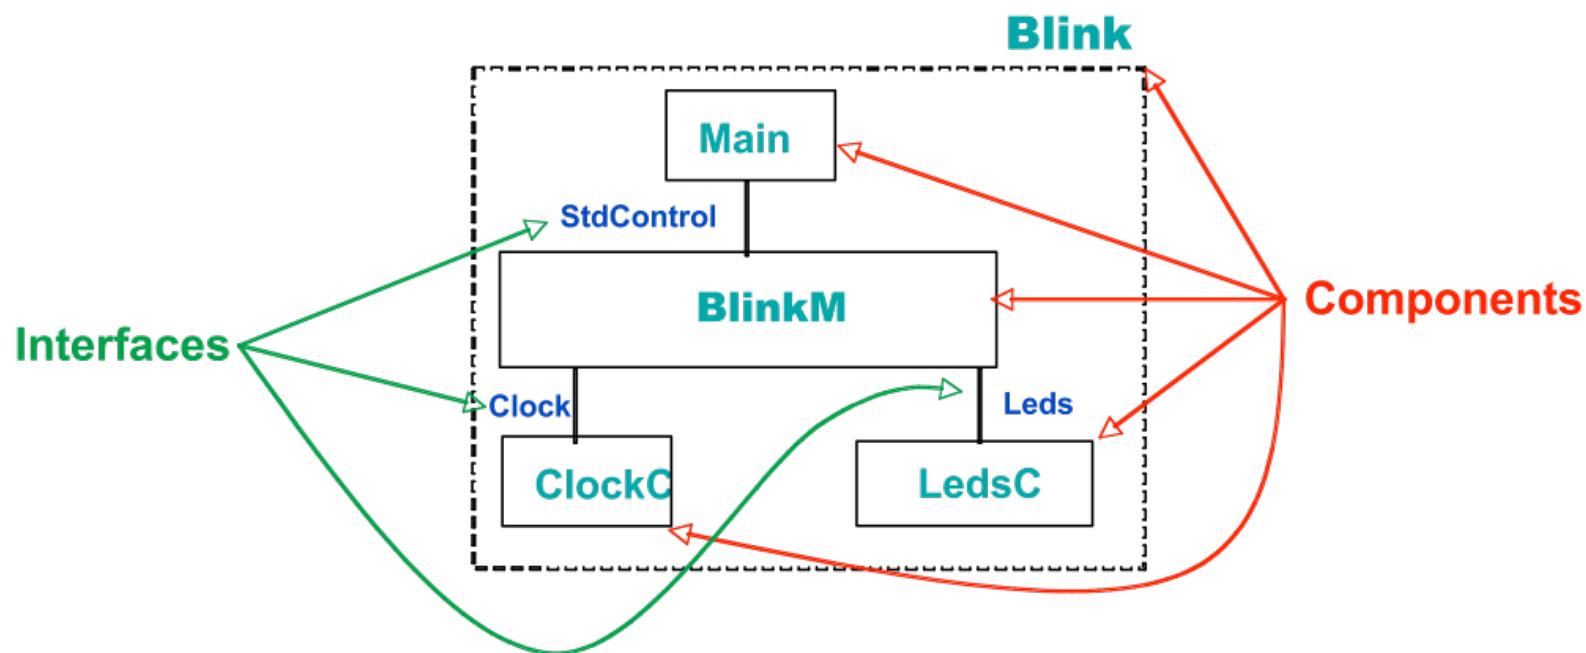
51

- Application consists of one or more components linked together to form an executable
- Two types of components
 - Modules: Implement functionality
 - Configurations: Just wire up other components
- Wiring (linking up) is done using interfaces
 - Every component provides and uses interfaces
 - Interfaces are the only points of access to a component
- Interfaces consist of two sets of functions
 - Commands: Providers of interfaces implement these functions
 - Events: Users of interfaces implement these functions
- A component may use multiple interfaces of the same or different type

Simple Application: Blink LEDs

52

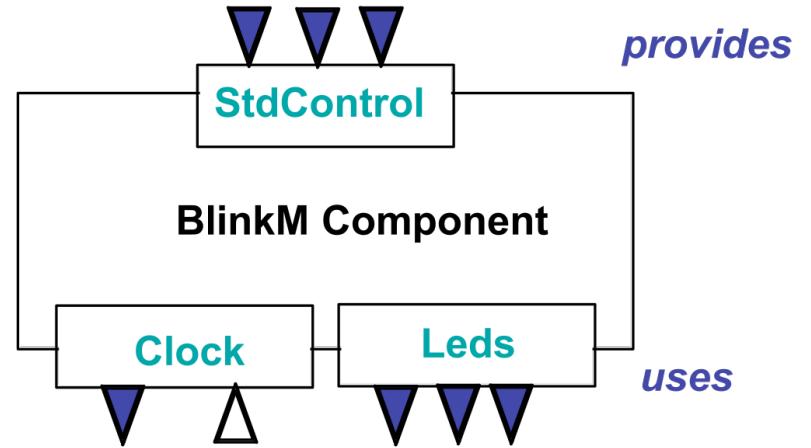
- Blink the red LED at 1Hz
- Application contains 5 components and 3 interfaces



Module Component

53

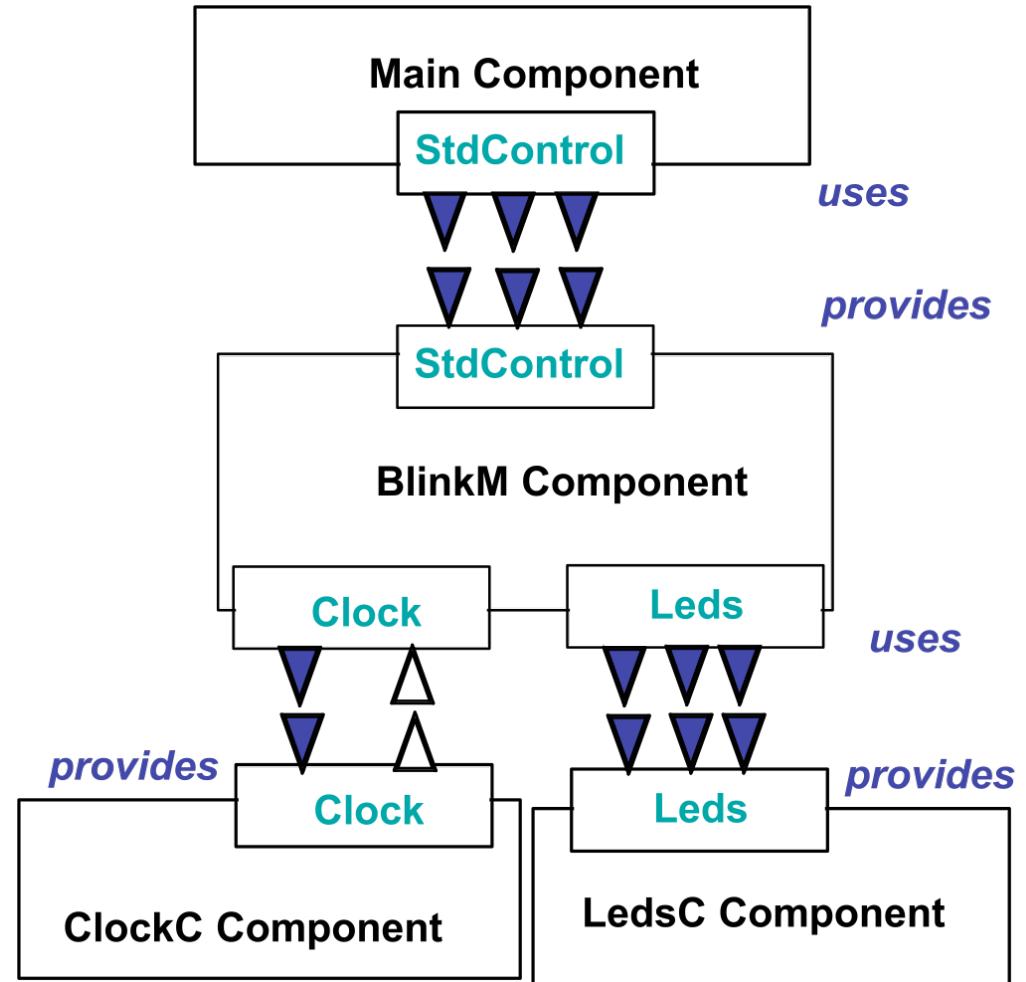
```
module BlinkM {  
    provides {  
        interface StdControl; }  
    uses {  
        interface Clock;  
        interface Leds; }  
    }  
  
    implementation {  
        bool state;  
        command result_t StdControl.init() {  
            state = FALSE;  
            return SUCCESS;  
        }  
        command result_t StdControl.start() {  
            return call Clock.setRate(TOS_I1PS, TOS_S1PS);  
        }  
        command result_t StdControl.stop() {  
            return call Clock.setRate(TOS_I0PS, TOS_S0PS);  
        }  
        event result_t Clock.fire() {  
            state = !state;  
            if (state)  
                call Leds.redOn();  
            else  
                call Leds.redOff();  
            return SUCCESS;  
        }  
    }
```



- Component is implementing functions
- Component is NOT wiring other components together
- Therefore, component is a module

The Complete Picture

```
Configuration Main {
    uses {
        interface StdControl;
    }
    module BlinkM {
        provides {
            interface StdControl;
        }
        uses {
            interface Clock;
            interface Leds;
        }
    }
}
Configuration ClockC {
    provides {
        interface Clock;
    }
}
Configuration LedsC {
    provides {
        interface Leds;
    }
}
```



TinyOS Concurrency Model

55

- Event-driven execution
 - Non-preemptive run-to-completion
- Two contexts
 - Foreground - Commands, Events
 - Background - Tasks
 - Background context can be interrupted by foreground

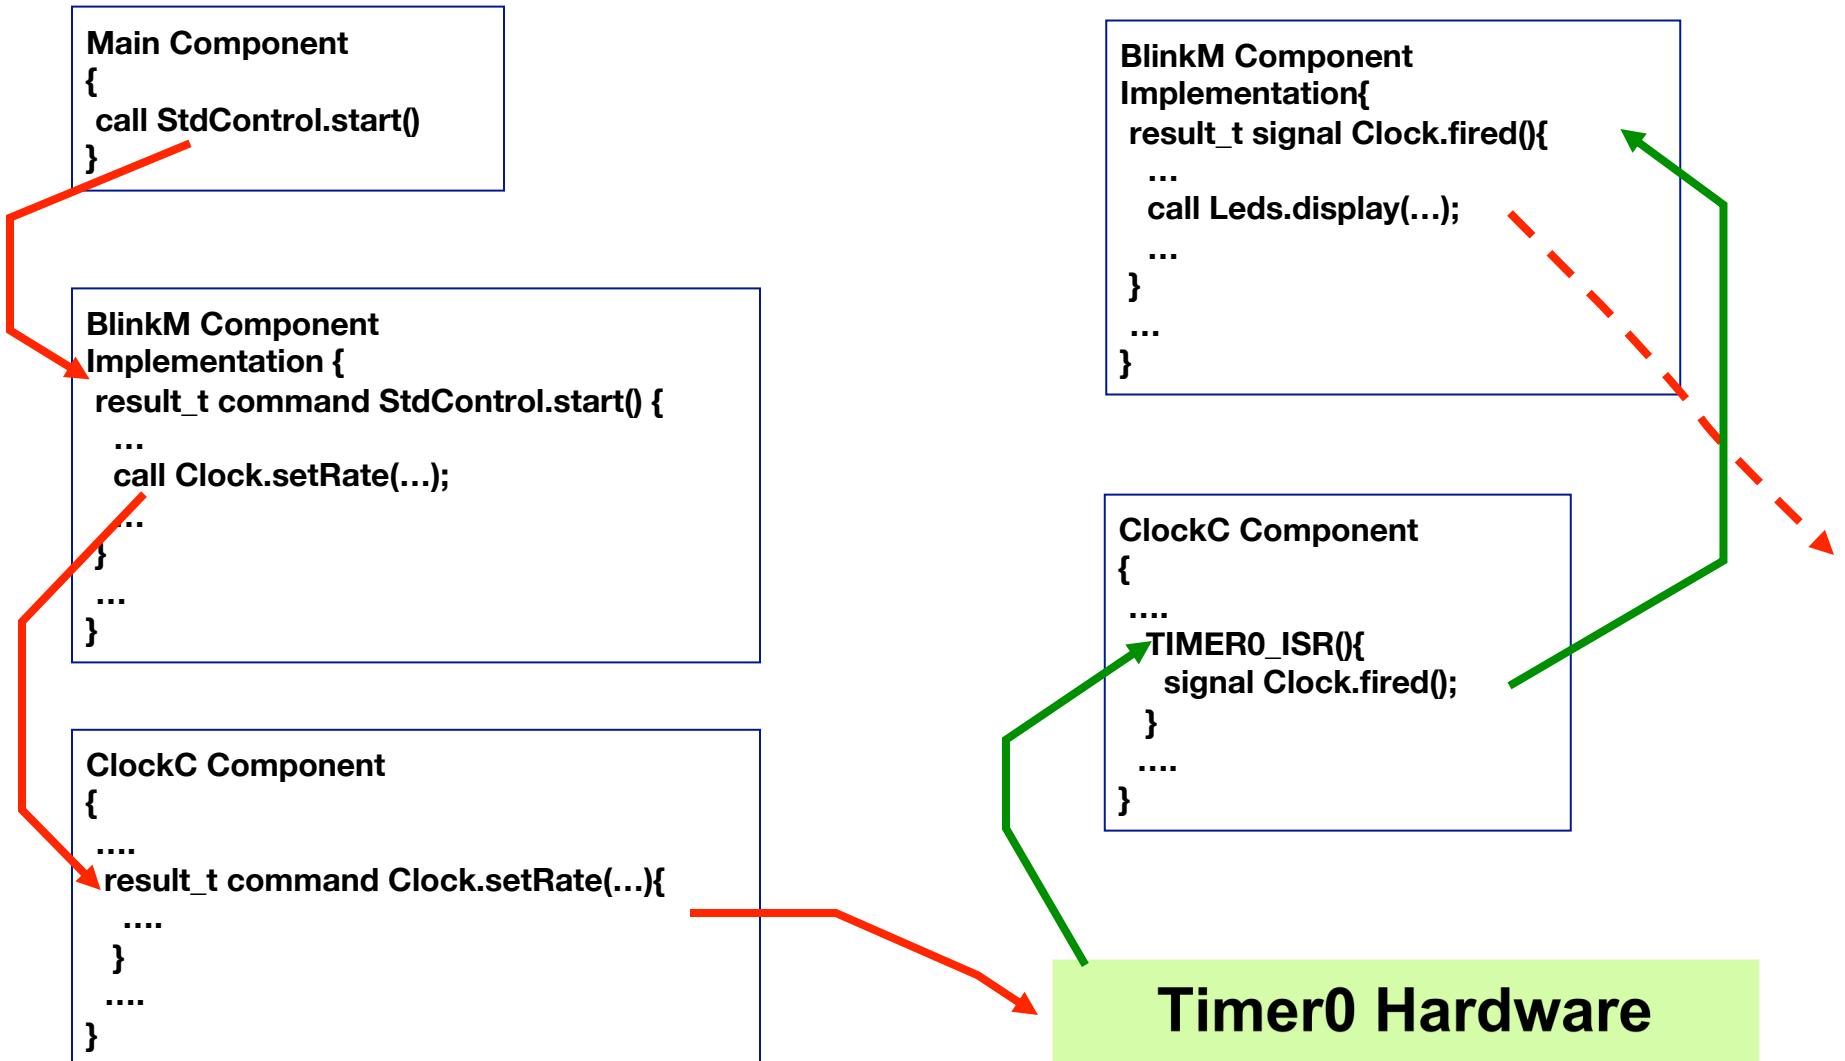
Clock.nc

```
interface Clock {  
    event result_t fire();  
    command result_t setRate(char interval,  
char scale);  
}
```

StdControl.nc

```
interface StdControl {  
    command result_t init();  
    command result_t start();  
    command result_t stop();  
}
```

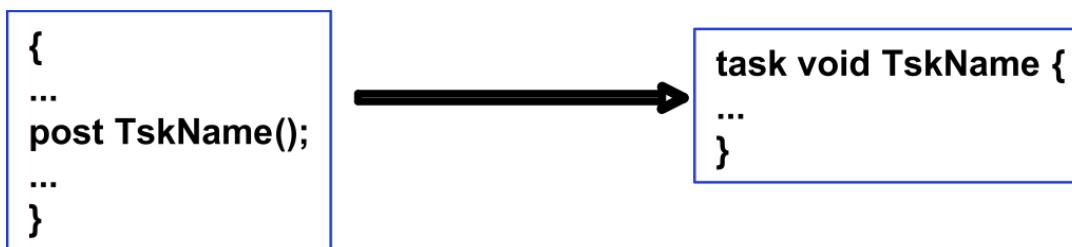
Understanding Blink Execution



Tasks

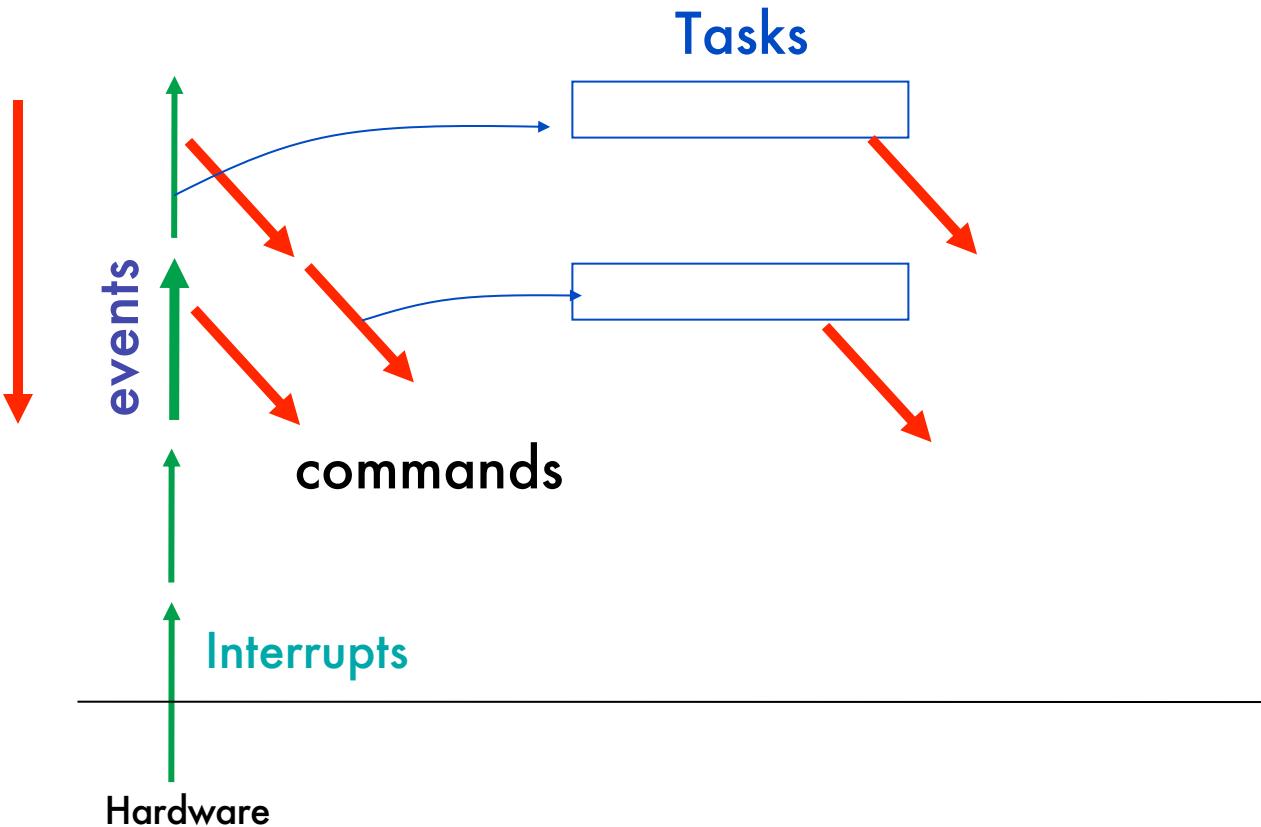
57

- To provide concurrency internal to a component
 - longer running operations (background)
- Can be preempted by events
- Able to perform operations beyond event context
 - e.g., process/transform data that is sampled
- May call commands
- May signal events
- Not preempted by other tasks
 - Run to completion model



TinyOS Execution Contexts

58



- Events generated by interrupts preempt tasks
- Tasks do not preempt tasks

TinyOS: Pros and Cons

59

- Pros
 - Static memory allocation => resource guarantee
 - Non-preemptive scheduling => minimal memory requirement
 - Modular nature => independent software development
- Cons
 - A new programming language
 - Static memory allocation => resource over subscription
 - Not really an OS
 - No resource management
 - Simple concurrency model
 - Monolithic firmware image hard to reconfigure dynamically

Example 2: μC/OS-II

60

- Portable, scalable, preemptive, multitasking RTOS
- Provides several services to applications
 - Semaphores, mutexes, event flags, mailboxes, message queues, task management, fixed-size memory block management, time management
- Source freely available for academic non-commercial usage for many platforms
 - Written almost entirely in C with some assembly
- In 2000, certified by Federal Aviation Administration (FAA) for use in avionics systems

µC/OS-II task model

61

- Users write applications in the form of tasks
- Each task is an infinite loop that runs on the CPU when “conditions” are satisfied and the kernel schedules it

```
void Task()
{
    //Initialisation code here

    while (1)
    {
        //Perform specific task here
    }
}
```

Task states

62

- Dormant
 - Exists in memory, but not known to kernel
- Ready
 - Task is in queue to use the CPU
- Running
 - Task has control of CPU and is executing
- Waiting
 - Task is suspended/blocked, waiting for something to happen

Simple Example

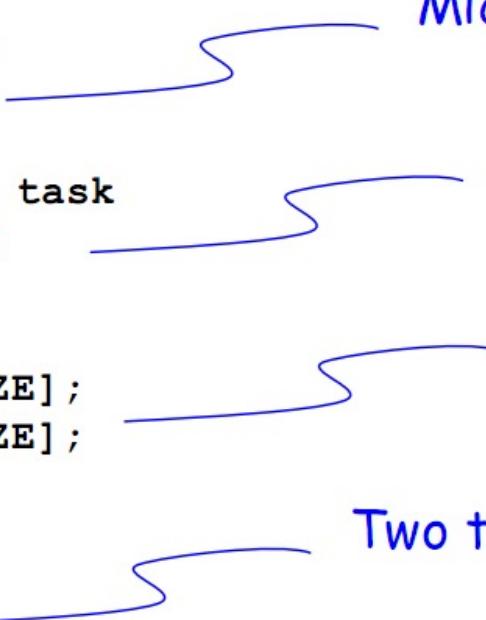
```
*****  
* uC/OS-II Win32 Example Program *  
*****  
  
// Include ucosII header file  
#include "includes.h"  
  
// Define stack size for each task  
#define TASK_STK_SIZE 512  
  
// Define Task stacks  
OS_STK Task1Stk[TASK_STK_SIZE];  
OS_STK Task2Stk[TASK_STK_SIZE];  
  
// Task Function Prototypes  
void Task1(void *pdata);  
void Task2(void *pdata);
```

MicroC/OS-II specific header

Stack size

Each task has a stack

Two task (function) prototypes



Simple Example

```
*****  
* First Task          *  
*****  
void Task1(void *pdata)  
{  
    printf("(II) Task1 initialised\n");  
    while (1) {  
        OSTimeDly(100);  
        printf("1 ");  
    }  
}  
  
*****  
* Second Task          *  
*****  
void Task2(void *pdata)  
{  
    printf("(II) Task2 initialised\n");  
    while (1)  
    {  
        OSTimeDlyHMSM(0,0,4,0);  
        printf("2 ");  
    }  
}
```

Normal function declaration

Enter infinite loop

Ask kernel to delay this task

Different delay function

Simple Example

```
int main()
{
    // Display a banner.
    printf("##### uCOS-II ELEC3730 Example 1\n");

    // Initialize uCOS-II.
    OSInit();                                Initialise the kernel (or OS)

    // Create some tasks
    OSTaskCreate(Task1, (void *) NULL, &Task1Stk[TASK_STK_SIZE], 5);
    OSTaskCreate(Task2, (void *) NULL, &Task2Stk[TASK_STK_SIZE], 6);

    // Start multitasking.
    OSSStart();                               Let kernel create TCB

    /* NEVER EXECUTED */
    printf("main(): We should never execute this line\n");
}
```

Normal old main routine

Initialise the kernel (or OS)

Let kernel create TCB

Start the kernel running

What happens after OSStart()?

66

- Kernel picks highest priority task and starts running it
 - Changes its state to “Running”
- When task asks kernel to delay its execution
 - Task goes into the “Waiting” state
 - Each clock “tick” event is captured by the kernel
 - Using a timer or clock ISR
 - At each tick, kernel decrements counter for each task and if the counter == 0, task state is changed to “Ready” state
 - If a task is highest priority in the “Ready” state, it gets selected for execution
 - Runs till the next time it asks kernel to delay it

Back to Example 1

67

```
int main()
{
    // Display a banner.
    printf("##### uCOS-II ELEC3730 Example 1\n");

    // Initialize uCOS-II.
    OSInit();

    // Create some tasks
    OSTaskCreate(Task1, (void *) NULL, &Task1Stk[TASK_STK_SIZE], 5);
    OSTaskCreate(Task2, (void *) NULL, &Task2Stk[TASK_STK_SIZE], 6);

    // Start multitasking.
    OSStart();
```

Kernel moves task state from dormant to ready

Start the kernel running moves highest priority ready task into running

Back to Example 1

68

```
*****  
* First Task          *  
*****  
void Task1(void *pdata)  
{  
    printf("(II) Task1 initialised\n");  
    while (1)  
    {  
        OSTimeDly(100);  
        printf("1 ");  
    }  
}  
  
*****  
* Second Task          *  
*****  
void Task2(void *pdata)  
{  
    printf("(II) Task2 initialised\n");  
    while (1)  
    {  
        OSTimeDlyHMSM(0,0,4,0);  
        printf("2 ");  
    }  
}
```

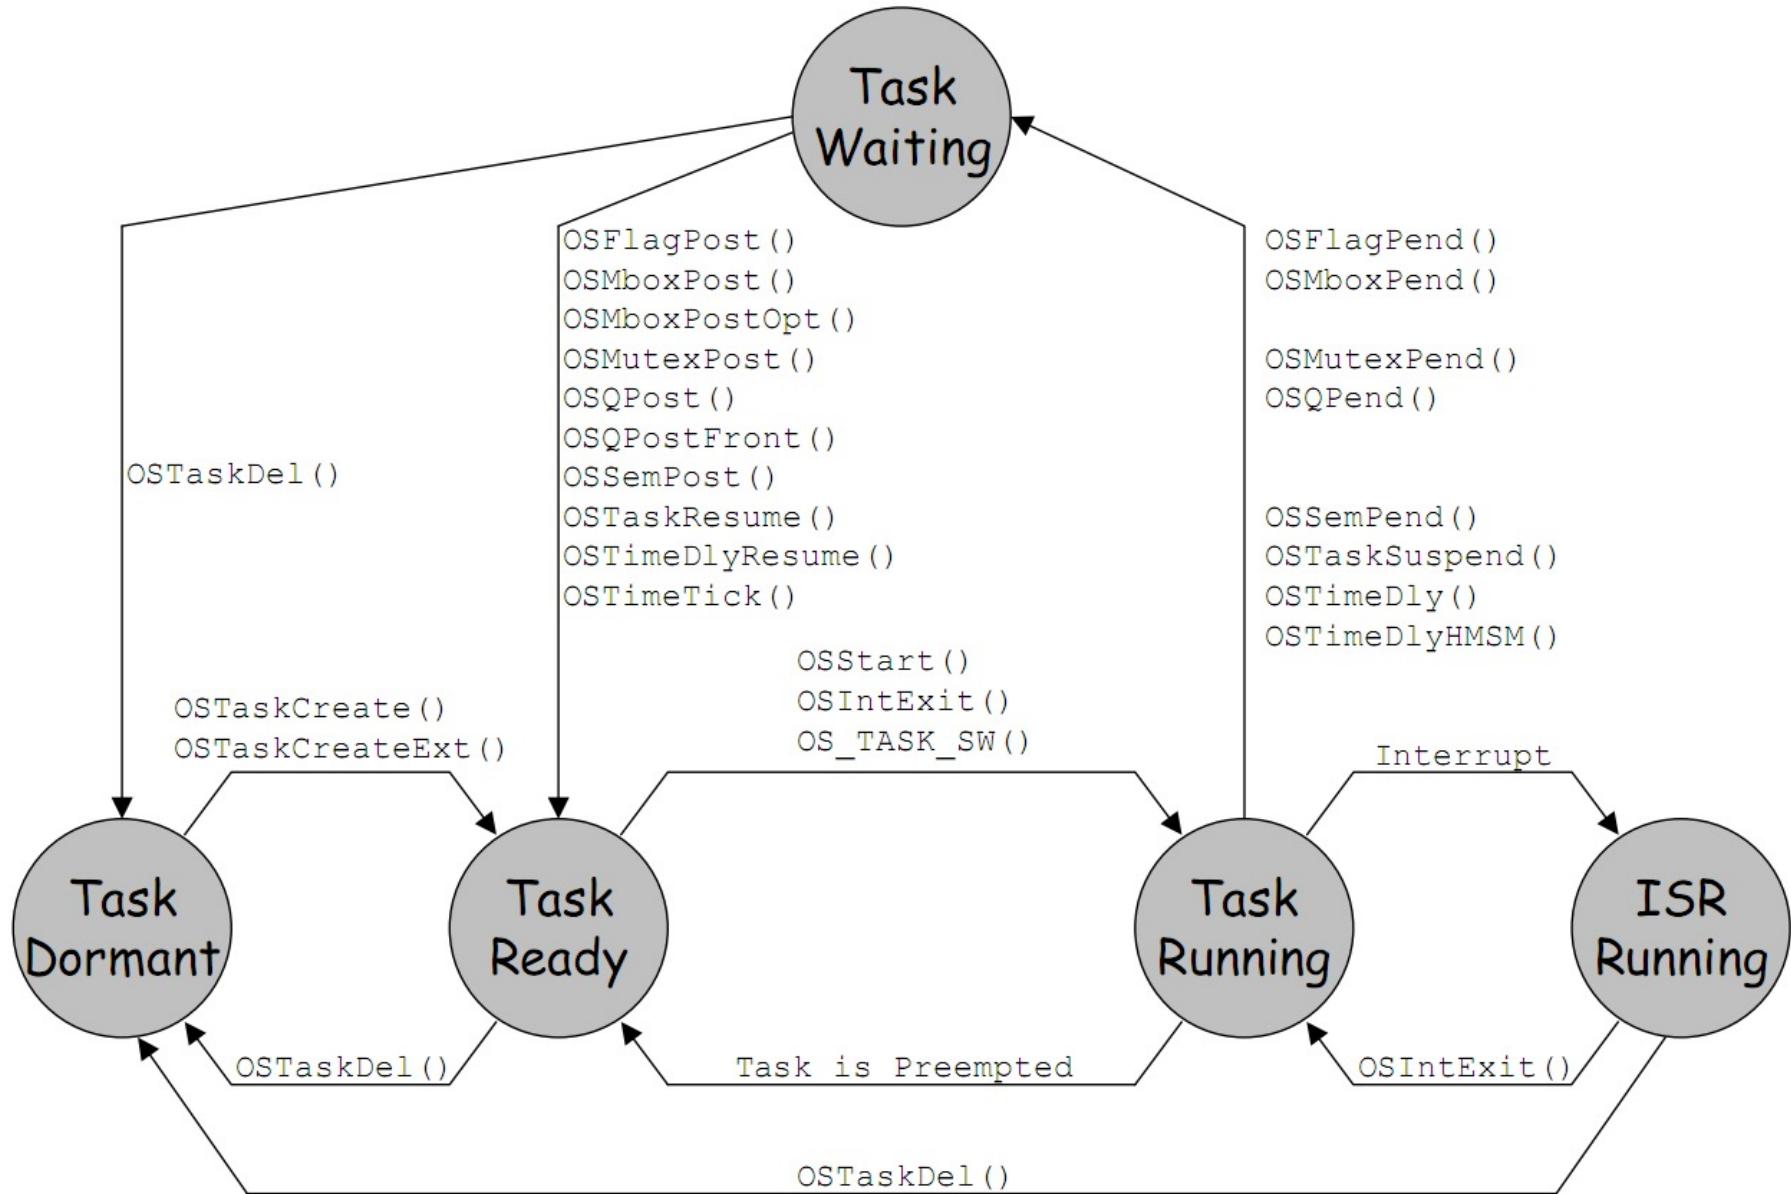
Kernel moves task from running to waiting state

Task 2 is now highest priority ready task, kernel moves from ready to running

Kernel moves task from running to waiting state

Tasks migrating between states

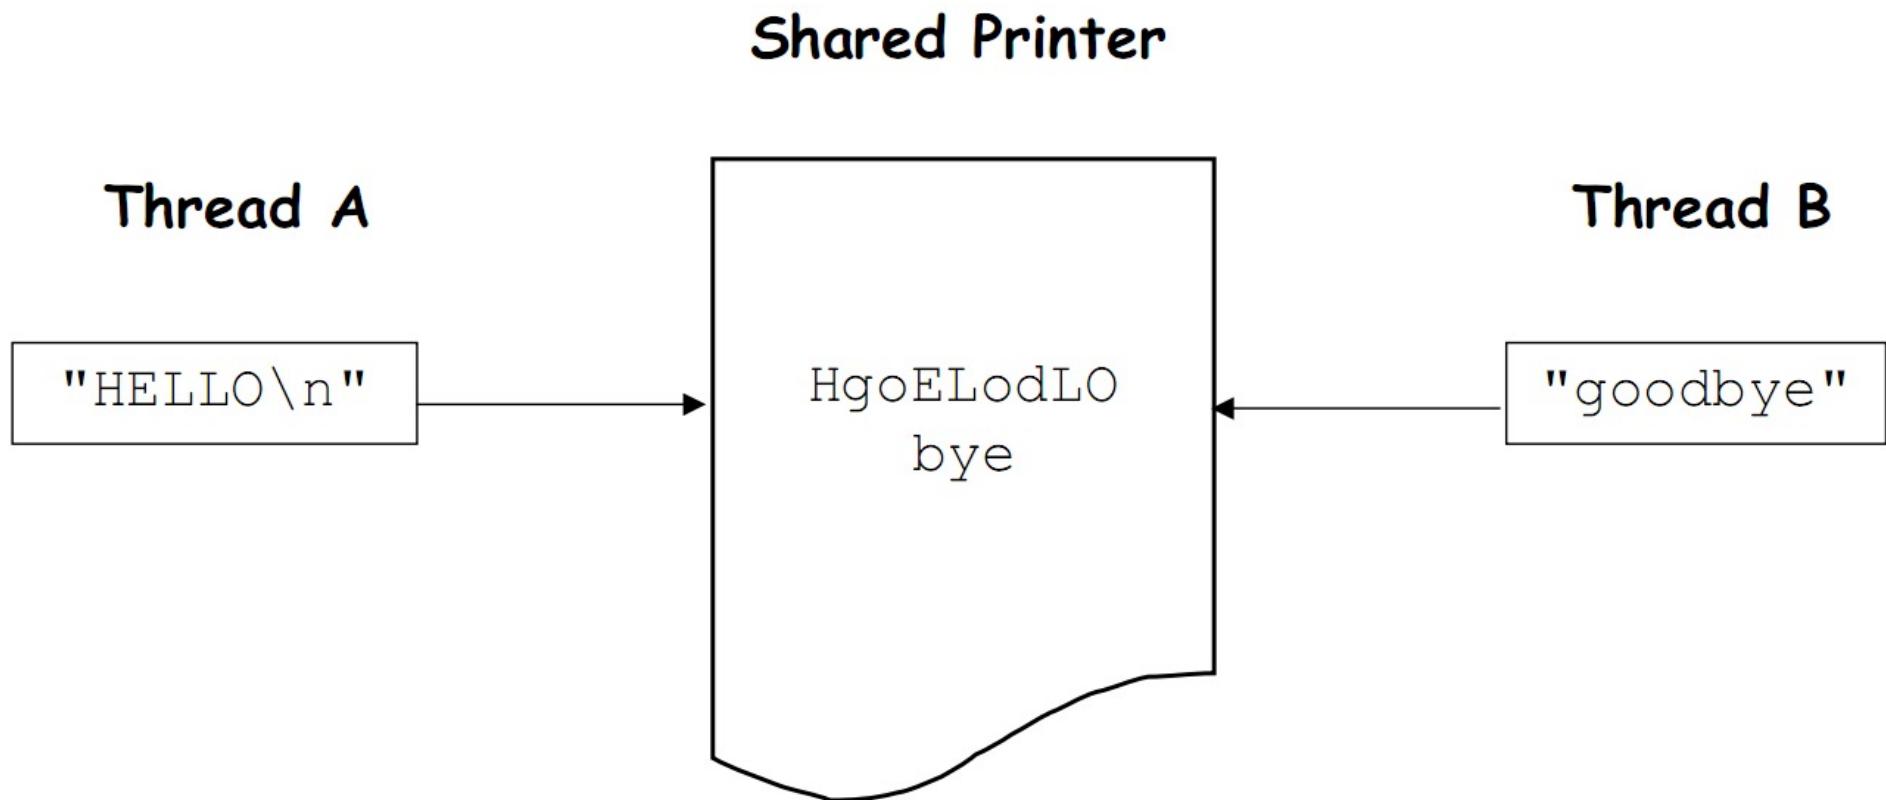
69



How about shared resources?

70

- Tasks often need exclusive access to shared resources
 - e.g., I/O devices, memory, network, etc.



Illustrating the problem in uC/OS-II

71

```
#include    "includes.h"
#define     TASK_STK_SIZE      512      // Stack size, in bytes
// Task stacks
OS_STK    Task1Stk[TASK_STK_SIZE];
OS_STK    Task2Stk[TASK_STK_SIZE];
// Task Function Prototypes
void Task1(void *pdata);
void Task2(void *pdata);

int SharedInteger;  Shared resource

int main()
{
    // Display a banner.
    printf("##### uCOS-II ELEC3730 Example 2\n");

    // Initialize uCOS-II.
    OSInit();

    // Create some tasks
    OSTaskCreate(Task1, (void *) NULL, &Task1Stk[TASK_STK_SIZE], 5);
    OSTaskCreate(Task2, (void *) NULL, &Task2Stk[TASK_STK_SIZE], 6);

    // Start multitasking.
    OSStart();

    /* NEVER EXECUTED */
    printf("main(): We should never execute this line\n");
}
```

Illustrating the problem in uC/OS-II

72

```
void Task1(void *pdata)
{
    printf("(II) Task1 initialised\n");
    while (1)
    {
        SharedInteger = 1; 
        OSTimeDly(50);
        printf("SharedInteger = %i\n", SharedInteger);
    }
}

void Task2(void *pdata)
{
    printf("(II) Task2 initialised\n");
    while (1)
    {
        OSTimeDly(75); 
        SharedInteger = 2; 
    }
}
```

Access to shared resource

Access to shared resource

Dealing with shared resources

73

- Three widely used methods to obtain exclusive access to a shared resource
 - Critical sections: Interrupts are disabled during these
 - Disable scheduling: Do not schedule any other tasks (ISRs can still run)
 - Semaphores: Tasks must obtain a special key/token before accessing the shared resource

Critical Sections

74

```
#define OS_ENTER_CRITICAL()    asm {PUSHF; CLI} /* Disable interrupts */  
  
#define OS_EXIT_CRITICAL()     asm POPF          /* Enable interrupts */  
  
int Temp  
  
void swap(int *x, int *y)  
{  
    OS_ENTER_CRITICAL();  
  
    Temp = *x;  
    *x = *y; ——————  
    *y = Temp;  
  
    OS_EXIT_CRITICAL();  
}
```

Interruptions disabled

Guaranteed execution order

Interruptions enabled

Disabling Scheduling

75

```
int Temp;  
  
void swap(int *x, int *y)  
{  
    OSSchedLock();  
  
    Temp = *x;  
    *x = *y; ————— Execution order, but interrupts enabled  
    *y = Temp;  
  
    OSSchedUnlock();  
}
```

Scheduler disabled

Execution order, but interrupts enabled

Scheduler enabled

- ISRs can still execute
- Useful if resource is not shared between tasks & ISRs
- Same task resumed after ISR even if higher priority tasks exist

Semaphores

77

- Invented by Edgar Dijkstra in 1960s
- Elegant mechanism to
 - Control access to a shared resource (mutual exclusion)
 - Signal the occurrence of an event
 - Allow two tasks to synchronize their activities
- Semaphore is a key (permit) that a task has to acquire in order to use a resource
- If some other task is using the key, a task will wait till the key is available (placed into the “Blocked” state)
- Two types of semaphores
 - Binary semaphores (0 or 1)
 - Counting semaphores (range depends on no. of bits used)

Semaphore Primitives

78

- Can perform three operations on a semaphore
 - Initialize (also called Create) – provides initial value
 - Wait (also called Pend)
 - Signal (also called Post)
- If a task wants semaphore, it will perform a Wait (Pend)
 - If value > 0, semaphore value is decremented and task runs
 - If value == 0, task is blocked
- Can specify a timeout
 - If semaphore not available within a timeout, move task to “Ready” state and signal an error to the task
- Task releases semaphore by calling Signal (Post)
 - If no task waiting on semaphore, increment value
 - If task waiting, give semaphore to waiting task

Semaphores in uC/OS-II

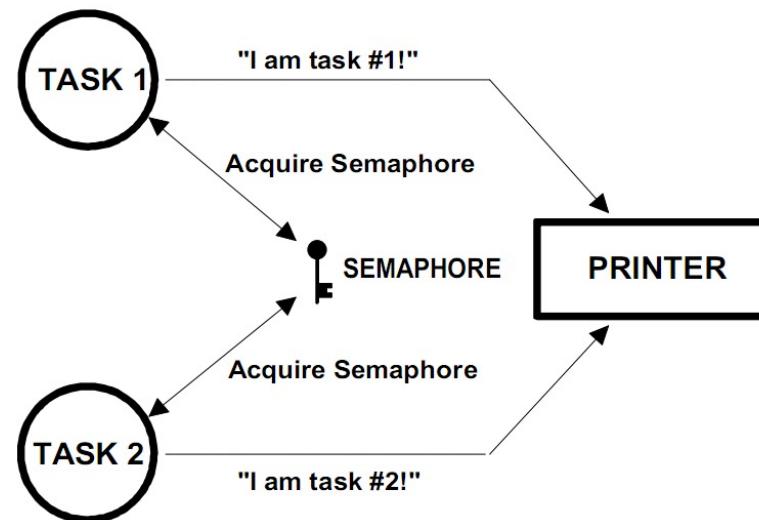
79

```
OS_EVENT *SharedDataSem;

void Function (void)
{
INT8U err;

OSSemPend(SharedDataSem, 0, &err);
.
/* You can access shared data in here (interrupts are recognized) */

OSSemPost(SharedDataSem);
}
```

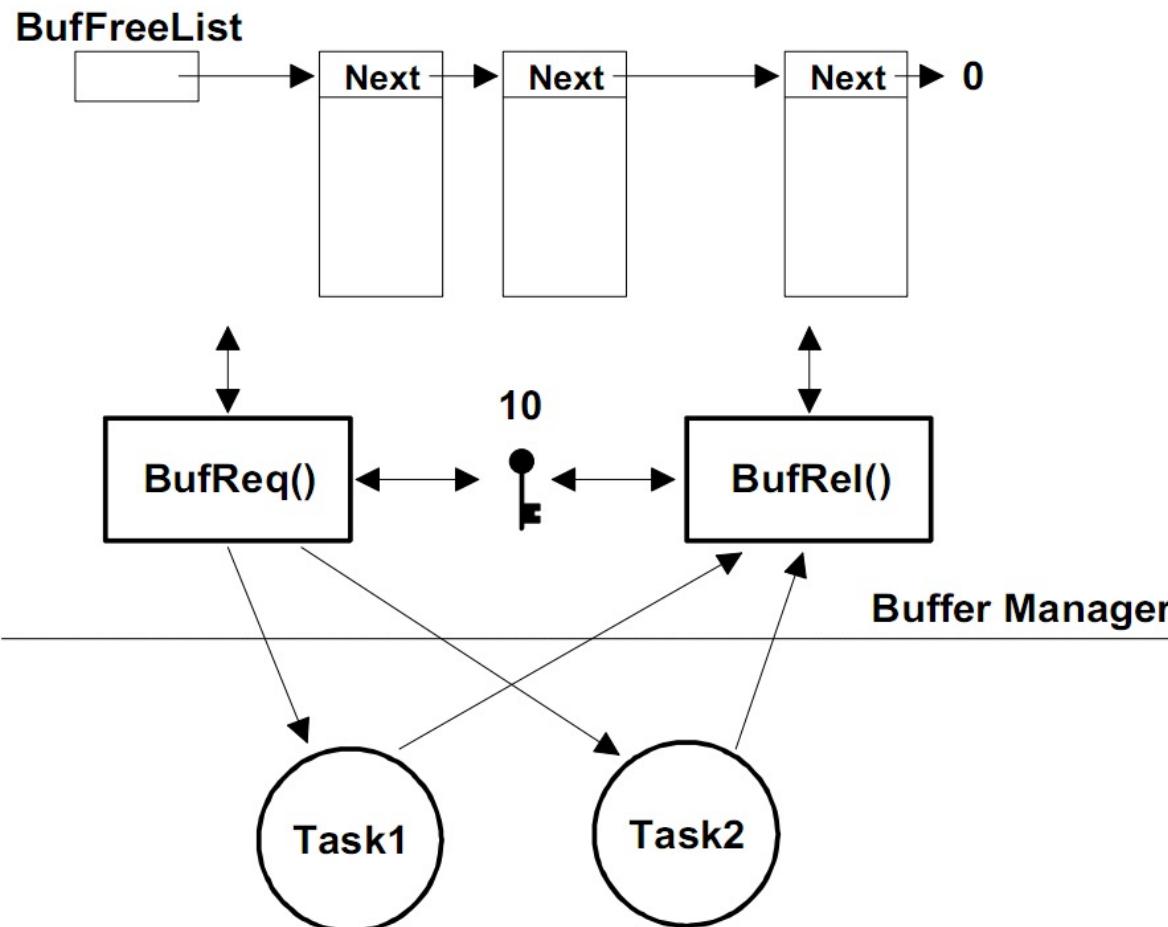


- Task needs to know about semaphore existence

Counting Semaphores Example

82

- Buffer management: You have a buffer pool (linked list) of 10 buffers that you want to give to tasks
 - ▶ Often used in networking protocols to store packets



```
BUF *BufReq(void)
{
    BUF *ptr;

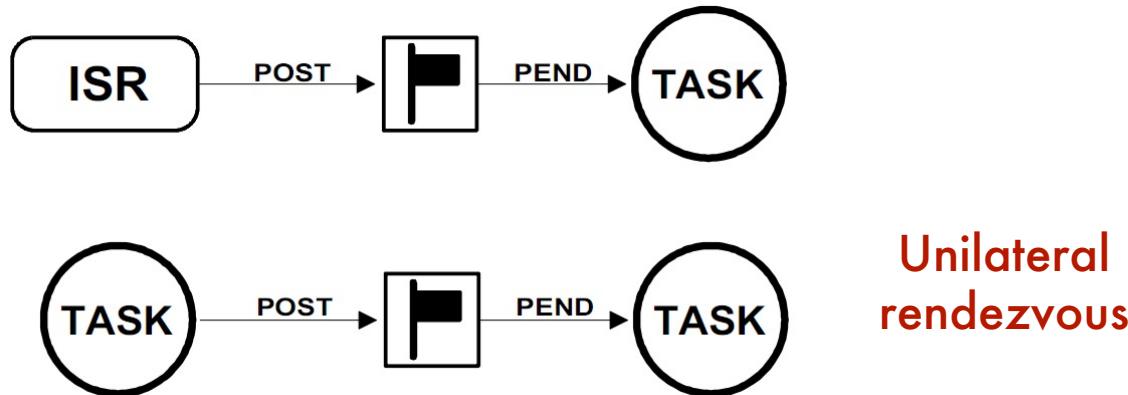
    Acquire a semaphore;
    Disable interrupts;
    ptr      = BufFreeList;
    BufFreeList = ptr->BufNext;
    Enable interrupts;
    return (ptr);
}

void BufRel(BUF *ptr)
{
    Disable interrupts;
    ptr->BufNext = BufFreeList;
    BufFreeList = ptr;
    Enable interrupts;
    Release semaphore;
}
```

Synchronization

83

- A task can be synchronized with an ISR or another task, using a semaphore



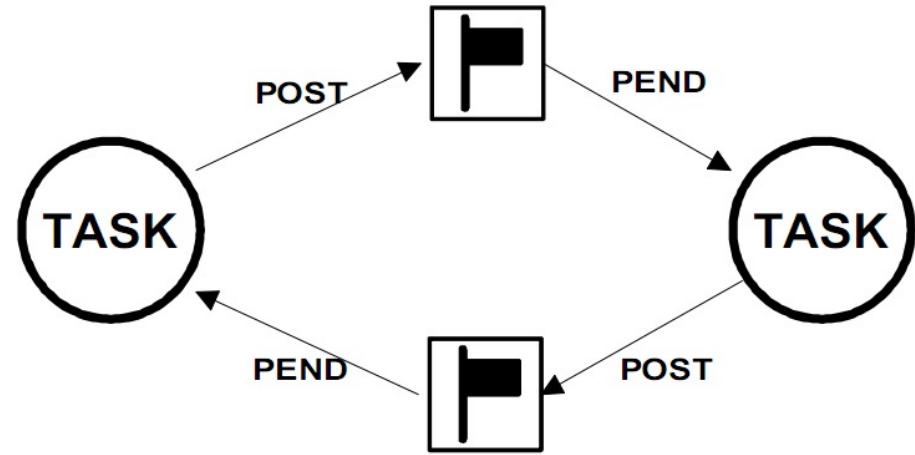
- Task initiates I/O to a peripheral (data packets to radio) and waits. An ISR or task signals the semaphore when I/O is complete and the task continues
- Initialize semaphore to 0, not 1
- More than one task may signal occurrence of event
- More than one task may be waiting for event

Bilateral Rendezvous

84

- Similar to unilateral rendezvous, but in both directions
 - Both tasks must synchronize with each other to continue
- Should not be used between task and ISRs

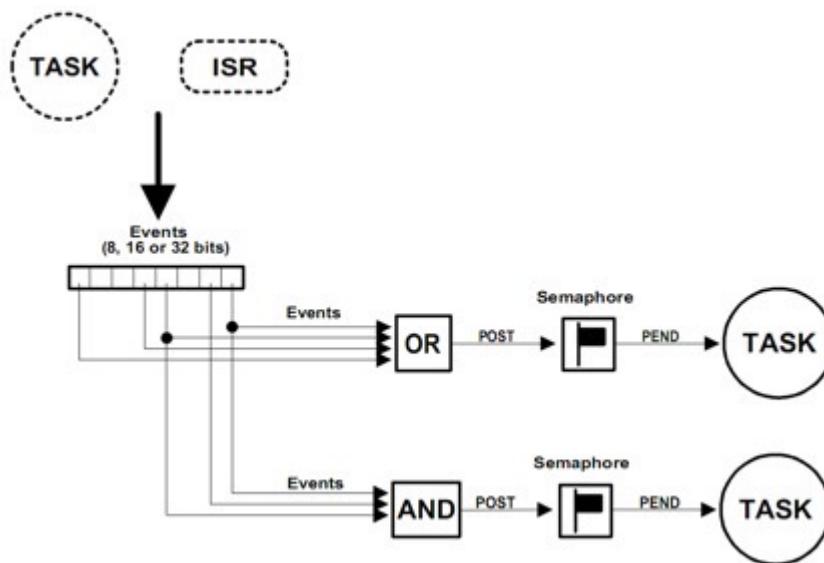
```
Task1 ()  
{  
    for (;;) {  
        Perform operation;  
        Signal task #2;  
        Wait for signal from task #2;  
        Continue operation;  
    }  
}  
  
Task2 ()  
{  
    for (;;) {  
        Perform operation;  
        Signal task #1;  
        Wait for signal from task #1;  
        Continue operation;  
    }  
}
```



Event Flags

85

- Task may be waiting for multiple events to take place
 - To synchronize with different events
- Disjunctive synchronization (logical OR)
 - When any of the events have taken place
- Conjunctive synchronization (logical AND)
 - When all of the events have taken place
- Tasks can set, clear, or wait for event flags



Intertask Communication

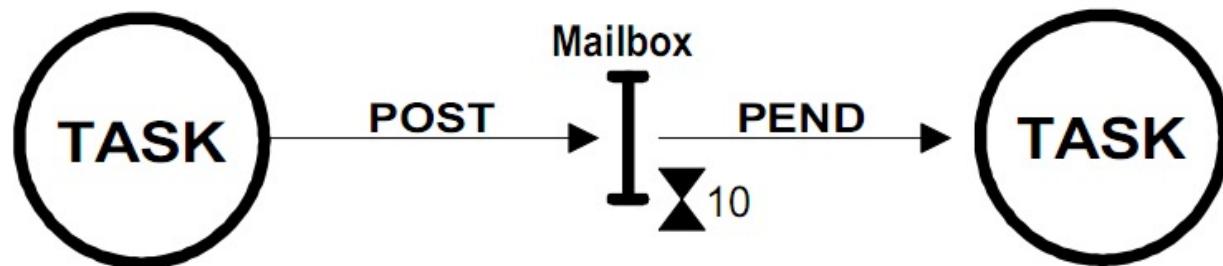
86

- Passing data (information) between ISRs/tasks can be done using global variables and data structures
 - Need to ensure mutual exclusion
 - ISR disabling or semaphore
- How does a task know that ISR has changed some global variable?
 - Option 1: Make sure ISR posts semaphore that task waits on
 - Option 2: Have task poll regularly to see if value changed
 - Option 3: Use a message mailbox or message queue

Message Mailboxes

87

- Usually, a pointer size variable
- Task/ISR can deposit a message (pointer) into mailbox
- One or more tasks can receive message from mailbox
 - Waiting list if multiple tasks wait for message in mailbox
 - Either priority or time used to arbitrate delivery
- Receivers of messages can specify timeouts
- Sender and receiver agree on what message represents
- Typical calls provided: Deposit, Wait, Accept



Reentrant Functions

89

- Reentrant function is one that can be called again before it completes
 - Multiple tasks can use function without data corruption
- Reentrant functions use only local variables or protect global variables when used

```
void strcpy(char *dest, char *src)
{
    while(*dest++ = *src++)
    {
        :
    }
    *dest = NULL;
}
```

Reentrant

```
int Temp

void swap(int *x, int *y)
{
    Temp = *x;
    *x = *y;
    *y = Temp;
}
```

Non-Reentrant

Reentrant Functions

90

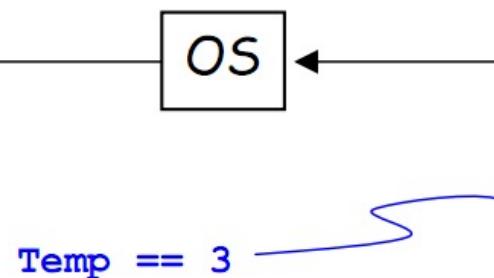
Low-priority task

```
while (1)
{
    int x = 1;
    int y = 2;

    swap(&x, &y)
    {
        Temp = *x;
        *x = *y;
        *y = Temp;
    }
    .
    .
    OSTimeDly(100);
}
```

Temp == 1

Context Switch



High-priority task

```
while (1)
{
    int z = 3;
    int t = 2;

    swap(&z, &t)
    {
        Temp = *z;
        *z = *t;
        *t = Temp;
    }
    .
    .
    OSTimeDly(100);
}
```

while (1)

{
int z = 3;
int t = 2;

swap(&z, &t)
{

Temp = *z;
*z = *t;
*t = Temp;

}

OSTimeDly(100);
.

Summary

91

- Software organization
 - Single program approach
 - Foreground background approach
 - Multitasking
 - Event driven non-preemptive
 - Preemptive multitasking
- Examples of RTOSs and standards
- Two RTOS internals: TinyOS and μ c/OS
- Dealing with shared resources, task communication