

Spring 2022

ECE 562: Embedded Systems

Lecture #6: Real-Time Scheduling (Fixed Priority Assignment)

Vijay Raghunathan

Modeling for Scheduling Analysis

2

- Need to specify workload model, resource model, and algorithm model of the system in order to analyze the timing properties
- Workload model: Describes the applications/tasks executed
 - Functional parameters: What does the task do?
 - Temporal parameters: Timing properties/requirements of the task
 - Precedence constraints and dependencies
- Resource model: Describes the resources available to the system
 - Number and type of CPUs, other shared resources
- Algorithm model: Describes how tasks use the available resources
 - Essentially a scheduling algorithm/policy

General Workload Model

3

- Set Γ of “n” tasks $\tau_1, \tau_2, \dots, \tau_n$ that provide some functionality
- Each task may be invoked multiple times as the system functions
 - Each invocation is referred to as a task instance
 - $\tau_{i,j}$ indicates the j^{th} instance of the i^{th} task
- Time when a task instance arrives (is activated/invoked) is called its release time (denoted by $r_{i,j}$)
- Each task instance has a run-time (denoted by $C_{i,j}$ or $e_{i,j}$)
 - May know range $[C_{i,j\min}, C_{i,j\max}]$
 - Can be estimated or measured by various mechanisms
- Each task instance has a deadline associated with it
 - Each task instance must finish before its deadline, else of no use to user
 - Relative deadline ($D_{i,j}$): Span from release time to when it must complete
 - Absolute deadline ($d_{i,j}$): Absolute wall clock time by which task instance must complete

Simplified Workload Model

4

- Tasks are periodic with constant inter-request intervals
 - Task periods are denoted by T_1, T_2, \dots, T_n
 - Request rate of τ_i is $1/T_i$
- Tasks are independent of each other
 - A task instance doesn't depend on the initiation/completion of other tasks
 - However, task periods may be related
- Execution time for a task is constant and does not vary with time
 - Can be interpreted as the maximum or worst-case execution time (WCET)
 - Denoted by C_1, C_2, \dots, C_n
- Relative deadline of every instance of a task is equal to the task period
 - $D_{i,j} = D_i = T_i$ for all instances $\tau_{i,j}$ of task τ_i
 - Each task instance must finish before the next request for it
 - Eliminates need for buffering to queue tasks
- Other implicit assumptions
 - No task can implicitly suspend itself, e.g., for I/O
 - All tasks are fully preemptible
 - All kernel overheads are zero

Resource Model

5

- Tasks need to be scheduled on one CPU
 - Referred to as uni-processor scheduling
 - Will relax this restriction later to deal with multi-processor scheduling
- Initially, will assume that there are no shared resources
 - Will relax this later to consider impact of shared resources on deadlines

Scheduling Algorithm

6

- Set of rules to determine the task to be executed at a particular moment
- One possibility: Preemptive and priority-driven
 - Tasks are assigned priorities
 - Statically or dynamically
 - At any instant, the highest priority task is executed
 - Whenever there is a request for a task that is of higher priority than the one currently being executed, the running task is preempted and the newly requested task is started
- Therefore, scheduling algorithm == method to assign priorities

Assigning Priorities to Tasks

7

- Static or fixed-priority approach
 - Priorities are assigned to tasks once
 - Every instance of the task has the same priority, determined apriori
- Dynamic approach
 - Priorities of tasks may change from instance to instance
- Mixed approach
 - Some tasks have fixed priorities, others don't

Deriving An Optimum Fixed Priority Assignment Rule

Critical Instant for a Task

9

- Overflow is said to occur at time t , if t is the deadline of an unfulfilled request
- A scheduling algorithm is feasible if tasks can be scheduled so that no overflow ever occurs
- Response time of a request of a certain task is the time span between the request and the end of response to that task
- Critical instant for a task = instant at which a request for that task will have the maximum response time
- Critical time zone of a task = time interval between a critical instant and the absolute deadline for that task instance

When does Critical Instant occur?

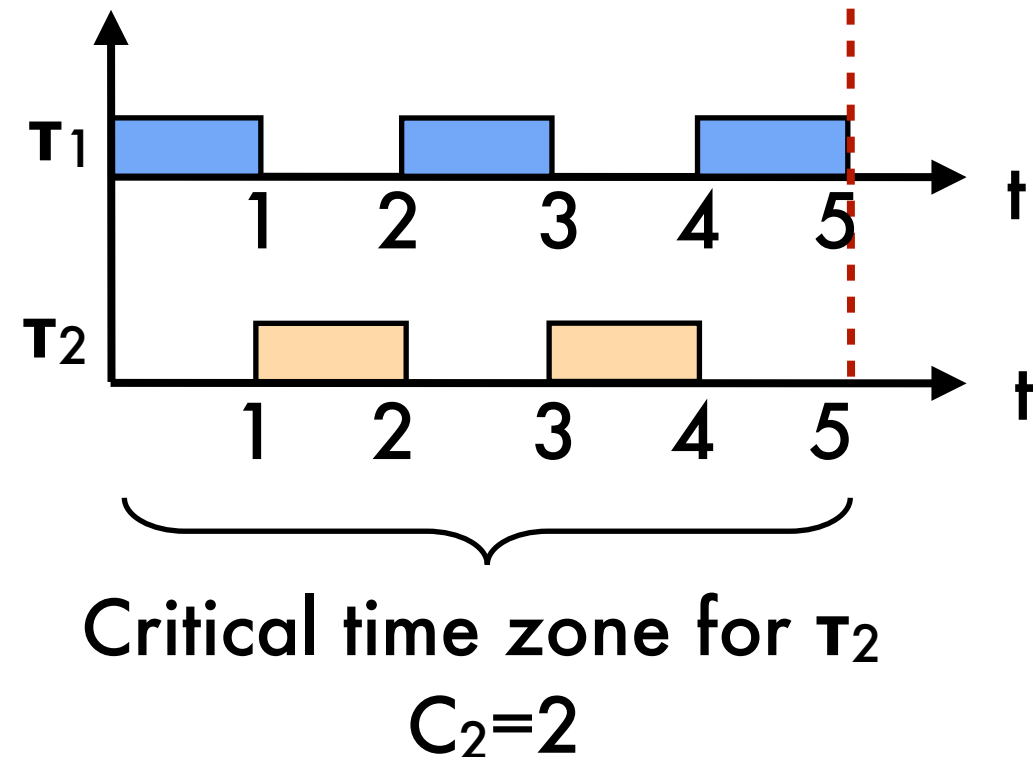
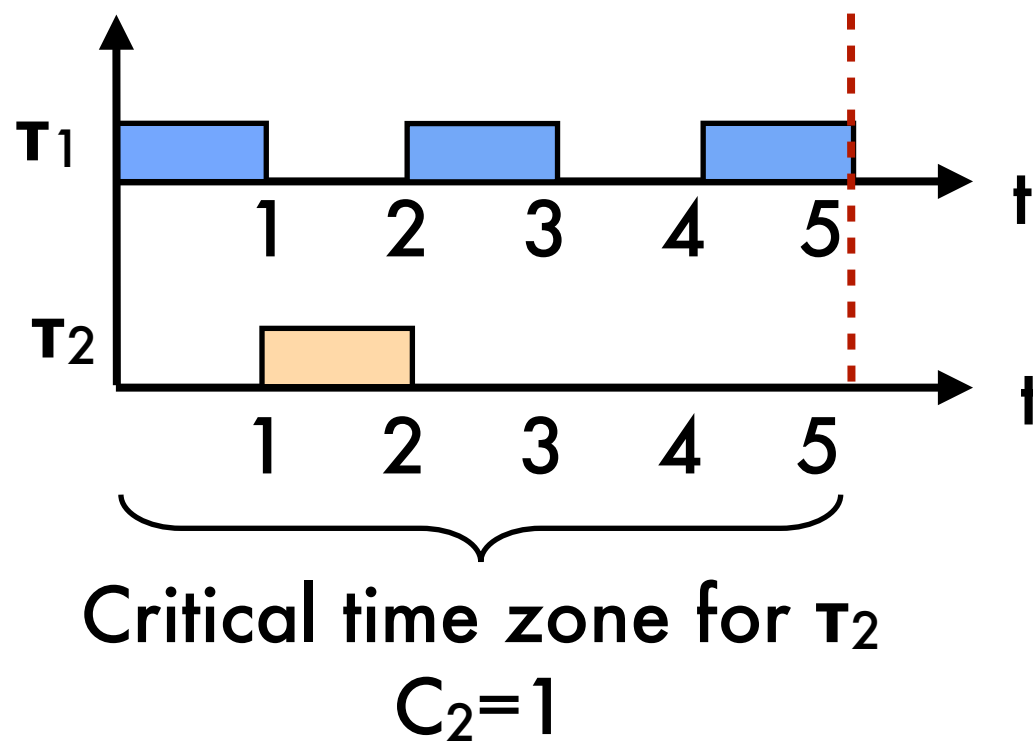
10

- **Theorem 1:** A critical instant for any task occurs whenever the task is requested simultaneously with requests of all higher priority tasks
- Can use this theorem to determine whether a given priority assignment will yield a feasible schedule or not
 - If requests for all tasks at their critical instants are fulfilled before their respective absolute deadlines, then the scheduling algorithm is feasible

Example

11


- Consider two tasks τ_1 and τ_2 with $T_1=2$, $T_2=5$, $C_1=1$, $C_2=1$
- Case 1: τ_1 has higher priority than τ_2
 - ▶ Priority assignment is feasible
 - ▶ Can increase C_2 to 2 and still avoid overflow



Example (contd.)

12

- Case 2: τ_2 has higher priority than τ_1
 - Priority assignment is still feasible
 - But, can't increase beyond $C_1=1, C_2=1$



Case 1 seems to be the better priority assignment for schedulability...
can we formalize this?

Rate-Monotonic Priority Assignment

13

- Assign priorities according to request rates, independent of execution times
 - Higher priorities for tasks with higher request rates (shorter time periods)
 - For tasks τ_i and τ_j , if $T_i < T_j$, $\text{Priority}(\tau_i) > \text{Priority}(\tau_j)$
- Called Rate-Monotonic (RM) Priority Assignment
 - It is optimal among static priority assignment based scheduling schemes
- **Theorem 2:** No other fixed priority assignment can schedule a task set if RM priority assignment can't schedule it, i.e., if a feasible priority assignment exists, then RM priority assignment is feasible

Proof of Theorem 2 (RM optimality)

14

- Consider n tasks $\{\tau_1, \tau_2, \dots, \tau_n\}$ ordered in increasing order of time periods (i.e., $T_1 < T_2 < \dots < T_n$)
- Assumption 1: Task set is schedulable with priority assignment $\{Pr(1), \dots, Pr(n)\}$ which is not RM
 - ▶ Therefore, \exists at least one pair of adjacent tasks, say τ_p and τ_{p+1} , such that $Pr(p) < Pr(p+1)$ [higher value is higher priority]
 - ▶ Otherwise, assignment becomes RM (violates assumption)
- Assumption 2: Instances of all tasks arrive at $t=0$
 - ▶ Therefore, $t=0$ is a critical instant for all tasks. From Theorem 1, we only need to check if first instance of each task completes before deadline

Proof (contd.)

15

- Swap the priorities of tasks τ_p and τ_{p+1}
 - New priority of τ_p is $Pr(p+1)$, new priority of τ_{p+1} is $Pr(p)$
 - Note that $Pr(p+1) > Pr(p)$ (by assumption 1)
- Tasks $\{\tau_1, \dots, \tau_{p-1}\}$ should not get affected
 - Since we are only changing lower priority tasks
- Tasks $\{\tau_{p+2}, \dots, \tau_n\}$ should also not get affected
 - Since both τ_p and τ_{p+1} need to be executed (irrespective of the order) before any task in $\{\tau_{p+2}, \dots, \tau_n\}$ gets executed
- Task τ_p should not get affected
 - Since we are only increasing its priority


Proof (contd.)

16

- Consider τ_{p+1} :
- Since original schedule is feasible, in the time interval $[0, T_p]$, exactly one instance of τ_p and τ_{p+1} complete execution along with (possibly multiple) instances of tasks in $\{\tau_1, \dots, \tau_{p-1}\}$
 - Note that τ_{p+1} executes before τ_p
- New schedule is identical, except that τ_p executes before τ_{p+1} (start/end times of higher priority tasks is same)
 - Still, exactly one instance of τ_p and τ_{p+1} complete in $[0, T_p]$. As $T_p < T_{p+1}$, task τ_{p+1} is schedulable

Proof (contd.)

17



We proved that swapping the priority of two adjacent tasks to make their priorities in accordance with RM does not affect the schedulability (i.e., all tasks $\{\tau_1, \tau_2, \dots, \tau_n\}$ are still schedulable)

Proof (contd.)

18

- If τ_p and τ_{p+1} are the only such non RM tasks in original schedule, we are done since the new schedule will be RM
- If not, starting from the original schedule, using a sequence of such re-orderings of adjacent task pairs, we can ultimately arrive at an RM schedule (Exactly the same as bubble sort)
- E.g., Four tasks with initial priorities $[3, 1, 4, 2]$ for $[\tau_1, \tau_2, \dots, \tau_n]$

$[3 \text{ } 1 \text{ } 4 \text{ } 2]$ is schedulable



$[3 \text{ } 4 \text{ } 1 \text{ } 2]$ is schedulable



$[4 \text{ } 3 \text{ } 1 \text{ } 2]$ is schedulable



$[4 \text{ } 3 \text{ } 2 \text{ } 1]$ is schedulable

Hence, Theorem 2 is proved.

RM priority
assignment

Processor Utilization Factor

19

- Processor Utilization: fraction of processor time spent in executing the task set (i.e., 1 - fraction of time processor is idle)
 - Provides a measure of computational load on CPU due to a task set
 - A task set is definitely not schedulable if its processor utilization is > 1
- For n tasks, $\tau_1, \tau_2, \dots, \tau_n$ the utilization “ U ” is given by:

$$U = C_1/T_1 + C_2/T_2 + \dots + C_n/T_n$$

- U for a task set Γ can be increased by increasing C_i 's or by decreasing T_i 's as long as tasks continue to satisfy their deadlines at their critical instants
- There exists a minimum value of U below which Γ is schedulable and above which it is not
 - Depends on scheduling algorithm and the task set


How Large can U be?

20

- For a given priority assignment, a task set **fully utilizes** a processor if:
 - ▶ The priority assignment is feasible for the task set (i.e., no deadline misses)
 - ▶ And, if an increase in the execution time of any task in the task set will make the priority assignment infeasible (i.e., cause a deadline miss)
- The U at which this happens is called the upper bound $U_{ub}(\Gamma, A)$ for a task set Γ under scheduling algorithm A
- The least upper bound of U is the minimum of the U's over all task sets that fully utilize the processor (i.e., $U_{lub}(A) = \min_{\Gamma} [U_{ub}(\Gamma, A)]$)
 - ▶ For all task sets whose U is below this bound, there exists a fixed priority assignment which is feasible
 - ▶ U above this can be achieved only if task periods T_i 's are suitably related
- $U_{lub}(A)$ is an important characteristic of a scheduling algorithm A as it allows easy verification of the schedulability of a task set
 - ▶ Below this bound, a task set is definitely schedulable
 - ▶ Above this bound, it may or may not be schedulable

Least Upper Bound of U

21



Utilization Bound for RM

22

- RM priority assignment is optimal; therefore, for a given task set, the U achieved by RM priority assignment is \geq the U for any other priority assignment
- In other words, the least upper bound of $U =$ the infimum of U 's for RM priority assignment over all possible T 's and all C 's for the tasks

Two Tasks Case

23

- **Theorem 3:** A set of two tasks is schedulable with fixed priority assignment if the processor utilization factor is $U \leq 2(2^{1/2}-1)$
- Given any task set consisting of only two tasks, if the utilization is less than 0.828, then the task set is schedulable using RM priority assignment

General Case

24

- **Theorem 4:** A set of n tasks with fixed priority assignment is schedulable if processor utilization factor $U \leq n(2^{1/n} - 1)$
- Equivalently, a set of “ n ” periodic tasks scheduled by the RM algorithm will always meet deadlines for all task start times if $C_1/T_1 + C_2/T_2 + \dots + C_n/T_n \leq n(2^{1/n} - 1)$

General Case (contd.)

25

- As $n \rightarrow \infty$, the U rapidly converges to $\ln 2 = 0.69$
- However, note that this is just the least upper bound
 - A task set with larger U may still be schedulable
 - e.g., if $(T_n \% T_i) = 0$ for $i=1,2,\dots,n-1$, then $U=1$
- How to check if a specific task set with n tasks is schedulable?
 - If $U \leq n(2^{1/n}-1)$ then it is schedulable
 - Otherwise, need to use Theorem 1!
- Two ways in which this analysis is useful
 - For a fixed CPU, will a set of tasks work or not? How much background load can you throw in without affecting feasibility of tasks?
 - During CPU design, you can decide how slow/fast a CPU you need

Theorem 1 Recalled

26

- Theorem 1: A critical instant for any task occurs whenever the task is requested simultaneously with requests of all higher priority tasks
- Can use this to determine whether a given priority assignment will yield a feasible scheduling algorithm
 - If requests for all tasks at their critical instants are fulfilled before their respective deadlines, then the scheduling algorithm is feasible
- Applicable to any static priority scheme... not just RM

Example #1

27

- Task τ_1 : $C_1 = 20$; $T_1 = 100$; $D_1 = 100$
Task τ_2 : $C_2 = 30$; $T_2 = 145$; $D_2 = 145$
Is this task set schedulable?
- $U = 20/100 + 30/145 = 0.41 \leq 2(2^{1/2}-1) = 0.828$
- Yes!

Example #2

28

- Task τ_1 : $C_1 = 20$; $T_1 = 100$; $D_1 = 100$
Task τ_2 : $C_2 = 30$; $T_2 = 145$; $D_2 = 145$
Task τ_3 : $C_3 = 68$; $T_3 = 150$; $D_3 = 150$
Is this task set schedulable?
- $U = 20/100 + 30/145 + 68/150 = 0.86 > 3(2^{1/3}-1) = 0.779$
- Can't say! Need to apply Theorem 1

Example #2 revisited

29

- The utilization based test is only a sufficient condition. Can we obtain a stronger test (a necessary and sufficient condition) for schedulability?
- Task τ_1 : $C_1 = 20$; $T_1 = 100$; $D_1 = 100$
Task τ_2 : $C_2 = 30$; $T_2 = 145$; $D_2 = 145$
Task τ_3 : $C_3 = 68$; $T_3 = 150$; $D_3 = 150$
- Consider the critical instant of τ_3 , the lowest priority task
 - ▶ τ_1 and τ_2 must execute at least once before τ_3 can begin executing
 - ▶ Therefore, completion time of τ_3 is $\geq C_1 + C_2 + C_3 = 20 + 68 + 30 = 118$
 - ▶ However, τ_1 is initiated one additional time in $(0, 118)$
 - ▶ Taking this into consideration, completion time of $\tau_3 = 2C_1 + C_2 + C_3 = 2 \cdot 20 + 68 + 30 = 138$
- Since $138 < D_3 = 150$, the task set is schedulable

Response Time Analysis for RM

30

- For the highest priority task, worst case response time R is its own computation time C
 - $R = C$
- Other lower priority tasks suffer interference from higher priority processes
 - $R_i = C_i + I_i$
 - I_i is the interference in the interval $[t, t+R_i]$

Response Time Analysis (contd.)

31

- Consider task i , and a higher priority task j
- Interference from task j during R_i :
 - ▶ # of releases of task j = $\lceil R_i/T_j \rceil$
 - ▶ Each release will consume C_j units of processor
 - ▶ Total interference from task j = $\lceil R_i/T_j \rceil * C_j$
- Let $hp(i)$ be the set of tasks with priorities higher than that of task i
- Total interference to task i from all tasks during R_i :

$$I_i = \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$$

Response Time Analysis (contd.)

32

- This leads to:

$$R_i = C_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$$

- Smallest R_i that satisfies the above equation will be the worst case response time
- Fixed point equation: can be solved iteratively

$$w_i^{n+1} = C_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n}{T_j} \right\rceil C_j$$

Algorithm

33

```
for i in 1..N loop -- for each process in turn
  n := 0
   $w_i^n := C_i$ 
  loop
    calculate new  $w_i^{n+1}$  from Equation
    if  $w_i^{n+1} = w_i^n$  then
       $R_i := w_i^n$ 
      exit {value found}
    end if
    if  $w_i^{n+1} > T_i$  then
      exit {value not found}
    end if
    n := n + 1
  end loop
end loop
```

Deadline Monotonic Assignment

34

- Relax the $D_i = T_i$ constraint to now consider $C_i \leq D_i \leq T_i$
- Priority of a task is inversely proportional to its relative deadline
 - $D_i < D_j \Rightarrow P_i > P_j$
- **DM is optimal**; Can schedule any task set that any other static priority assignment can
- Example: RM fails but DM succeeds for the following task set

	Period T	Deadline D	Comp Time, C	Priority P	Response Time, R
Task_1	20	5	3	4	3
Task_2	15	7	3	3	6
Task_3	10	10	4	2	10
Task_4	20	20	3	1	20

- Schedulability Analysis: One approach is to reduce task periods to relative deadlines
 - $C_1/D_1 + C_2/D_2 + \dots + C_n/D_n \leq n(2^{1/n}-1)$
 - However, this is very pessimistic
- A better approach is to do critical instant (response time) analysis

Task Synchronization

35

- So far, we considered independent tasks
- In reality, tasks do interact: semaphores, locks, monitors, rendezvous, etc.
 - shared data, use of non-preemptable resources
- Jeopardizes systems ability to meet timing constraints
 - e.g., may lead to an indefinite period of “priority inversion” where a high priority task is prevented from executing by a low priority task

Priority Inversion Example

36

- Let τ_1 and τ_3 be two tasks that share a resource (protected by semaphore S), with τ_1 having a higher priority. Let τ_2 be an intermediate priority task that does not share any resource with either. Consider the following sequence of actions:
- τ_3 gets activated, obtains a lock on the semaphore S , and starts using the shared resource
- τ_1 becomes ready to run and preempts τ_3 . While executing, τ_1 tries to use the shared resource by trying to lock S . But S is already locked and therefore τ_1 is blocked
- Now, τ_2 becomes ready to run. Since only τ_2 and τ_3 are ready to run, τ_2 preempts τ_3 .

Priority Inversion Example (contd.)



37

- What would we prefer?
 - ▶ τ_1 , being the highest priority task, should be blocked no longer than the time τ_3 takes to complete its critical section
- But, in reality, the duration of blocking is unpredictable
 - ▶ τ_3 can remain preempted until τ_2 (and any other pending intermediate priority tasks) are completed
- The duration of priority inversion becomes a function of the task execution times, and is not bounded by the duration of critical sections

Just another theoretical problem?

38

- Recall the Mars Pathfinder from 1997?
 - ▶ Unconventional landing - bouncing onto Martian surface with airbags
 - ▶ Deploying the Sojourner rover: First roving probe on another planet
 - ▶ Gathering and transmitting voluminous data, including panoramic pictures that were such a hit: http://en.wikipedia.org/wiki/Mars_Pathfinder
 - ▶ Used VxWorks real-time kernel (preemptive, static-priority scheduling)
- But...
 - ▶ A few days into the mission, not long after Pathfinder started gathering meteorological data, the spacecraft began experiencing total system resets, each resulting in losses of data
 - ▶ Reported in the press as "software glitches" and "the computer was trying to do too many things at once"



What really happened on Mars?

39

- The failure was a priority inversion failure!
- A high priority task *bc_dist* was blocked by a much lower priority task *ASI/MET* which had grabbed a shared resource and was then preempted by a medium priority communications task
- The high priority *bc_dist* task didn't finish in time
- An even higher priority scheduling task, *bc_sched*, periodically creates transactions for the next bus cycle
- *bc_sched* checks whether *bc_dist* finished execution (hard deadline), and if not, resets the system

Why was it not caught before launch?

40

- The problem only manifested itself when *ASI/MET* data was being collected and intermediate tasks were heavily loaded
- Before launch, testing was limited to the "best case" data rates and science activities
- Did see the problem before launch but could not get it to repeat when they tried to track it down
 - Neither reproducible or explainable
 - Attributed to "hardware glitches"
 - Lower priority - focus was on the entry and landing software

What saved the day?

41

- How did they find the problem?
 - ▶ Trace/log facility + a replica on earth
- How did they fix it?
 - ▶ Changed the creation flags for the semaphore so as to enable “priority inheritance”
 - ▶ VxWorks supplies global configuration variables for parameters, such as the “options” parameter for the `semMCreate` used by the select service
 - Turns out that the Pathfinder code was such that this global change worked with minimal performance impact
 - ▶ **Spacecraft code was patched: sent “diff”**
 - Custom software on the spacecraft (with a whole bunch of validation) modified the onboard copy

Diagnosing the Problem

42

- Diagnosing the problem as a black box would have been impossible
- Only detailed traces of actual system behavior enabled the faulty execution sequence to be captured and identified
- See http://research.microsoft.com/en-us/um/people/mbj/mars_pathfinder/ for a description of how things were diagnosed and fixed

Process Interactions and Blocking

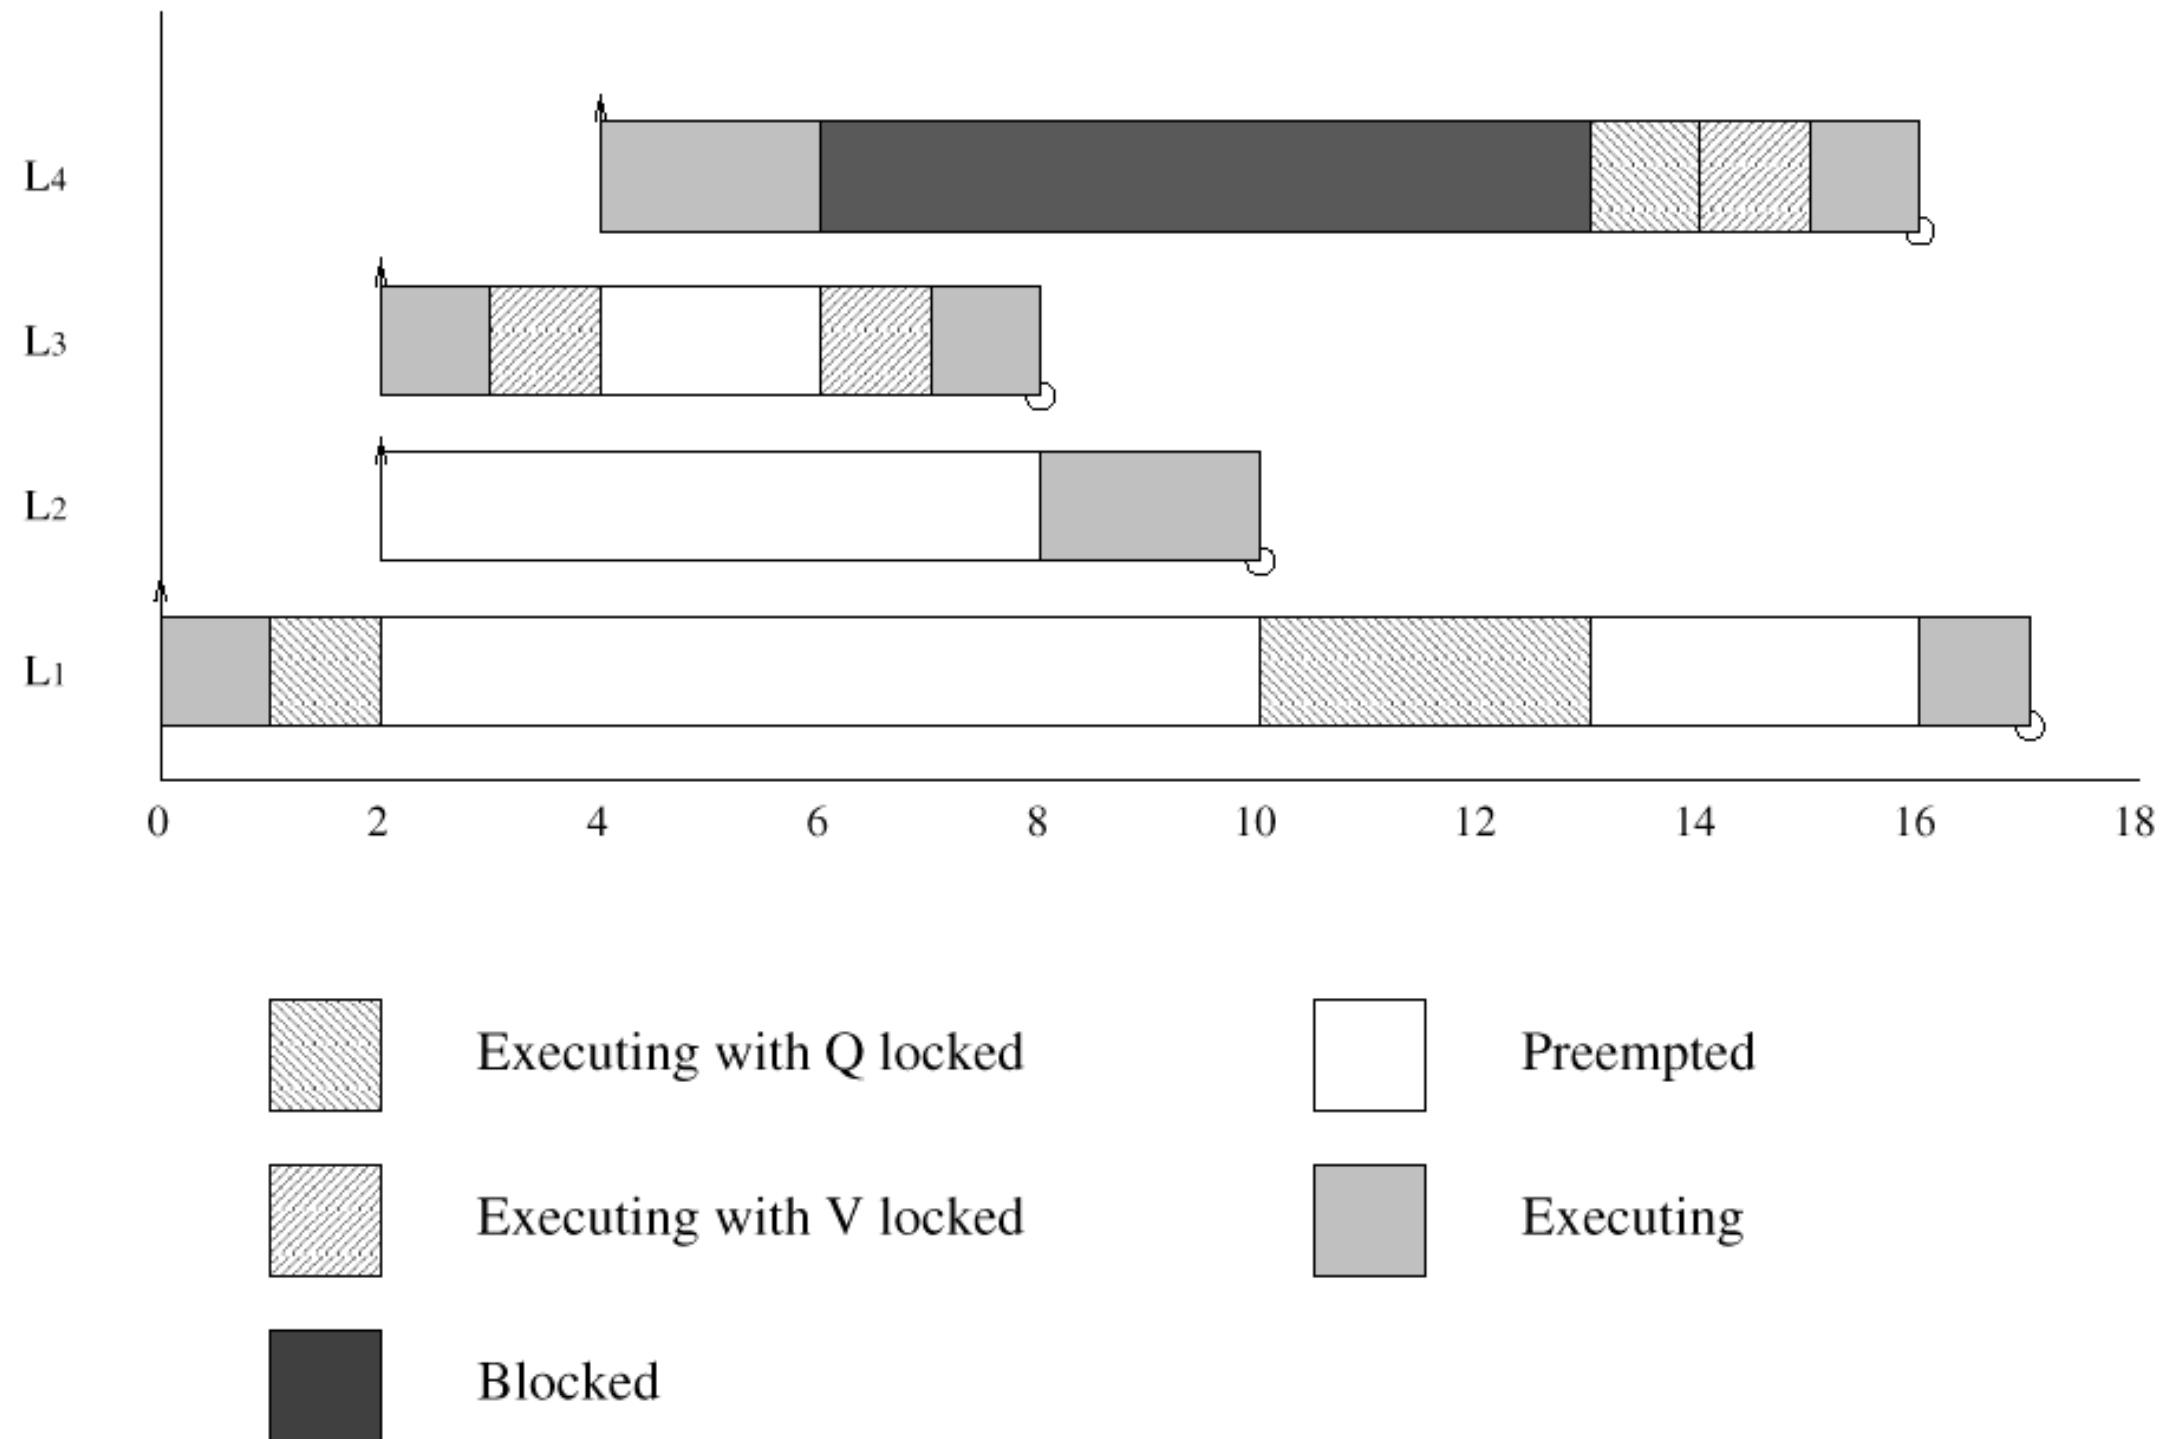
43

- Priority inversions
- Blocking
- Priority inheritance

Process	Priority	Execution Seq	Release Time
L_4	4	EEQVE	4
L_3	3	EVVE	2
L_2	2	EE	2
L_1	1	EQQQQE	0

Example: Priority Inversion

44




Priority Inheritance

45

- Simple method for eliminating priority inversion problems
- **Basic Idea:** If a high priority task H gets blocked while trying to lock a semaphore that has already been locked by a low priority task L, then L temporarily *inherits* the priority of H while it holds the lock to the semaphore
 - ▶ The moment L releases the semaphore lock, its priority drops back down
- Any intermediate priority task, I, will not preempt L because L will now be executing with a higher priority while holding the lock

Example: Priority Inheritance

46



Response Time Calculations

47

- $R = C + B + I$
 - solve by forming recurrence relation
- With priority inheritance:

$$R_i = C_i + B_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$$

$$B_i = \sum_{k=1}^K usage(k, i) CS(k)$$

Response Time Calculations

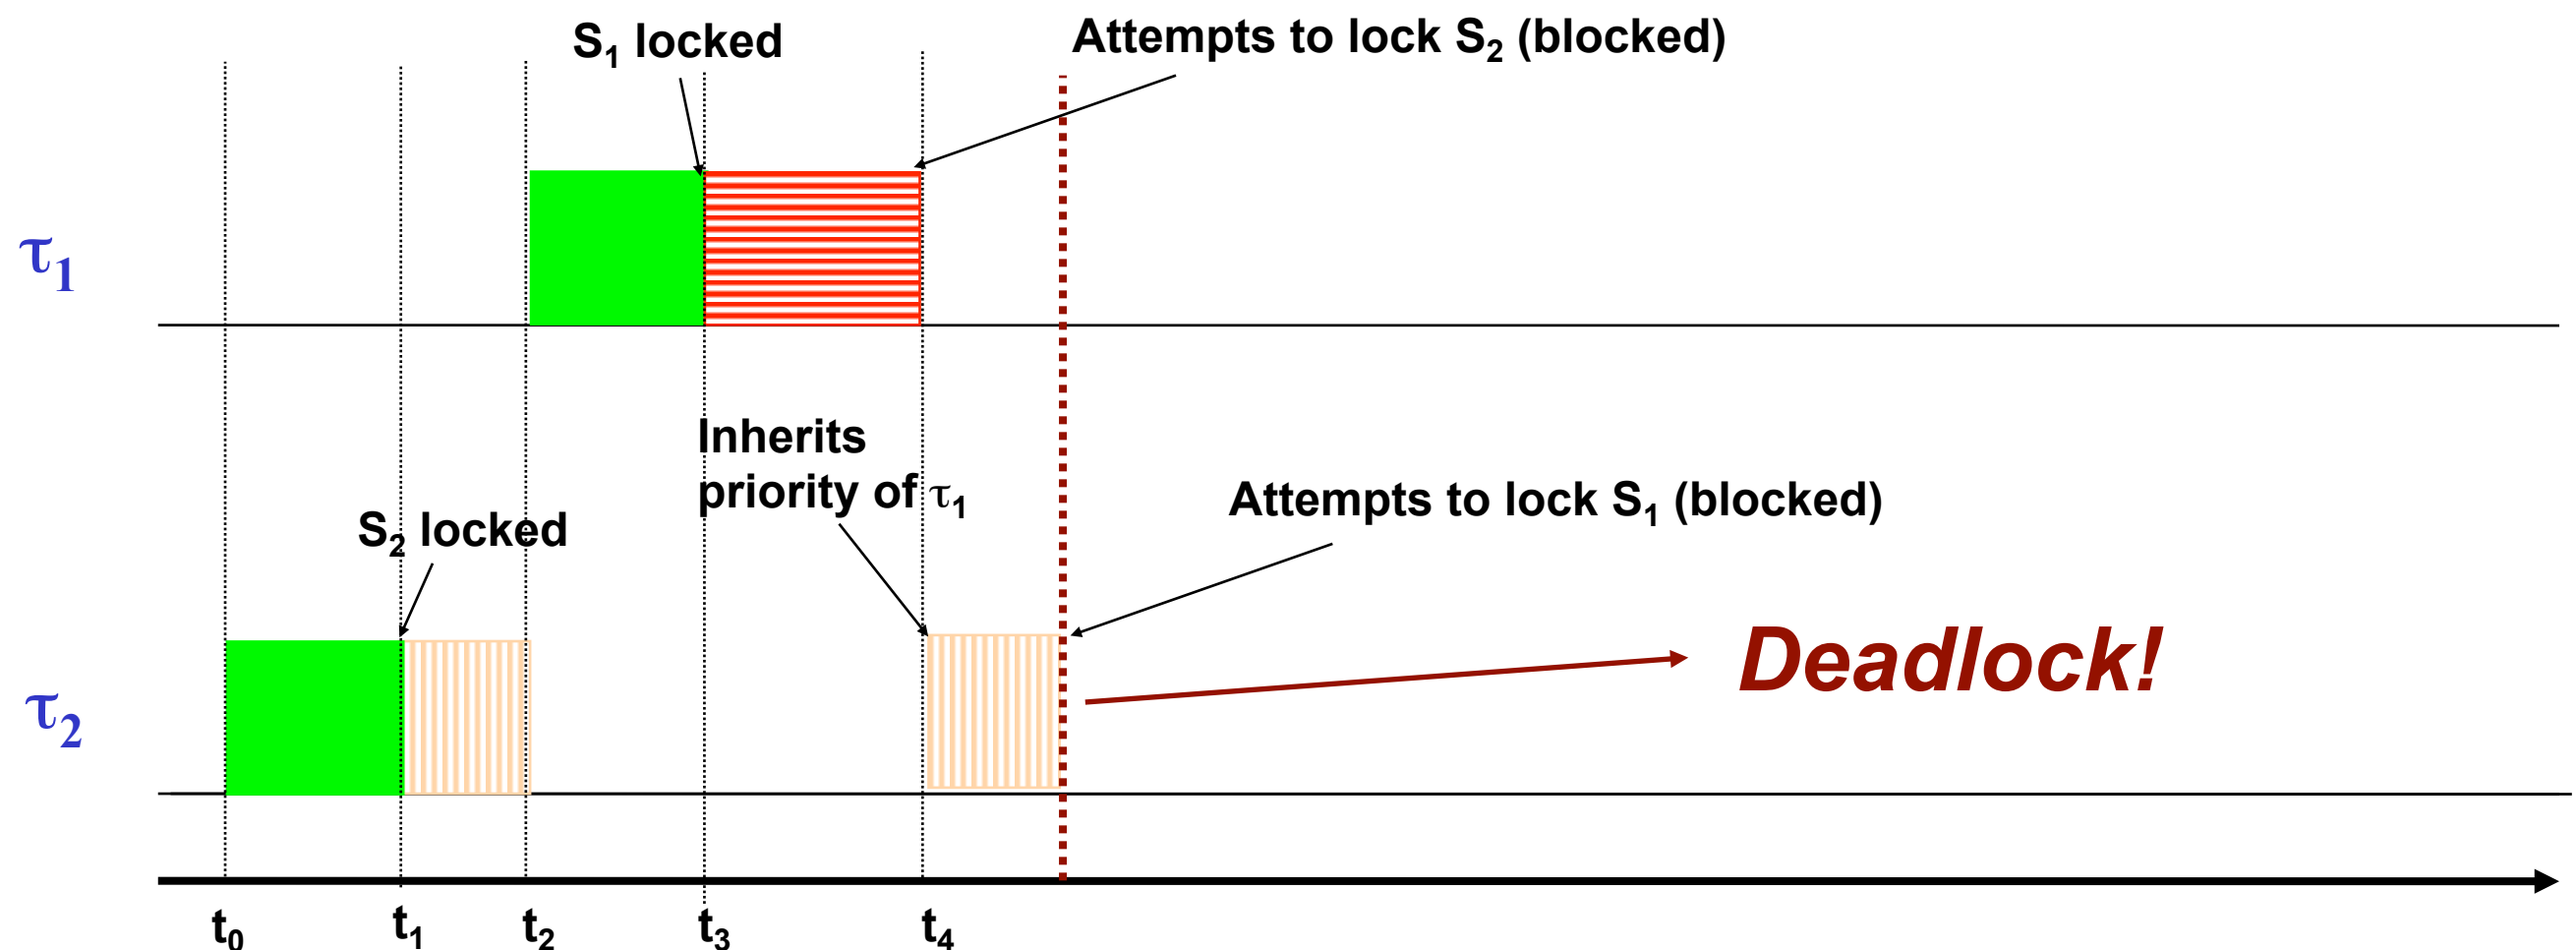
48

- Where usage is a 0/1 function:
 - ▶ $\text{usage}(k, i) = 1$ if resource k is used by at least 1 process with priority $< i$,
and at least one process with a priority greater or equal to i .
 - ▶ $= 0$ otherwise
- $\text{CS}(k)$ is the computational cost of executing the critical section associated with resource k

Priority Inheritance Can Lead to Deadlock

49

- Two tasks τ_1 and τ_2 with two shared data structures protected by binary semaphores S_1 and S_2 .
 - τ_1 : { ... Lock(S_1) ... Lock(S_2) ... Unlock(S_2) ... Unlock(S_1) ... }
 - τ_2 : { ... Lock(S_2) ... Lock(S_1) ... Unlock(S_1) ... Unlock(S_2) ... }
- Assume τ_1 has higher priority than τ_2



Priority Ceiling Protocols

50

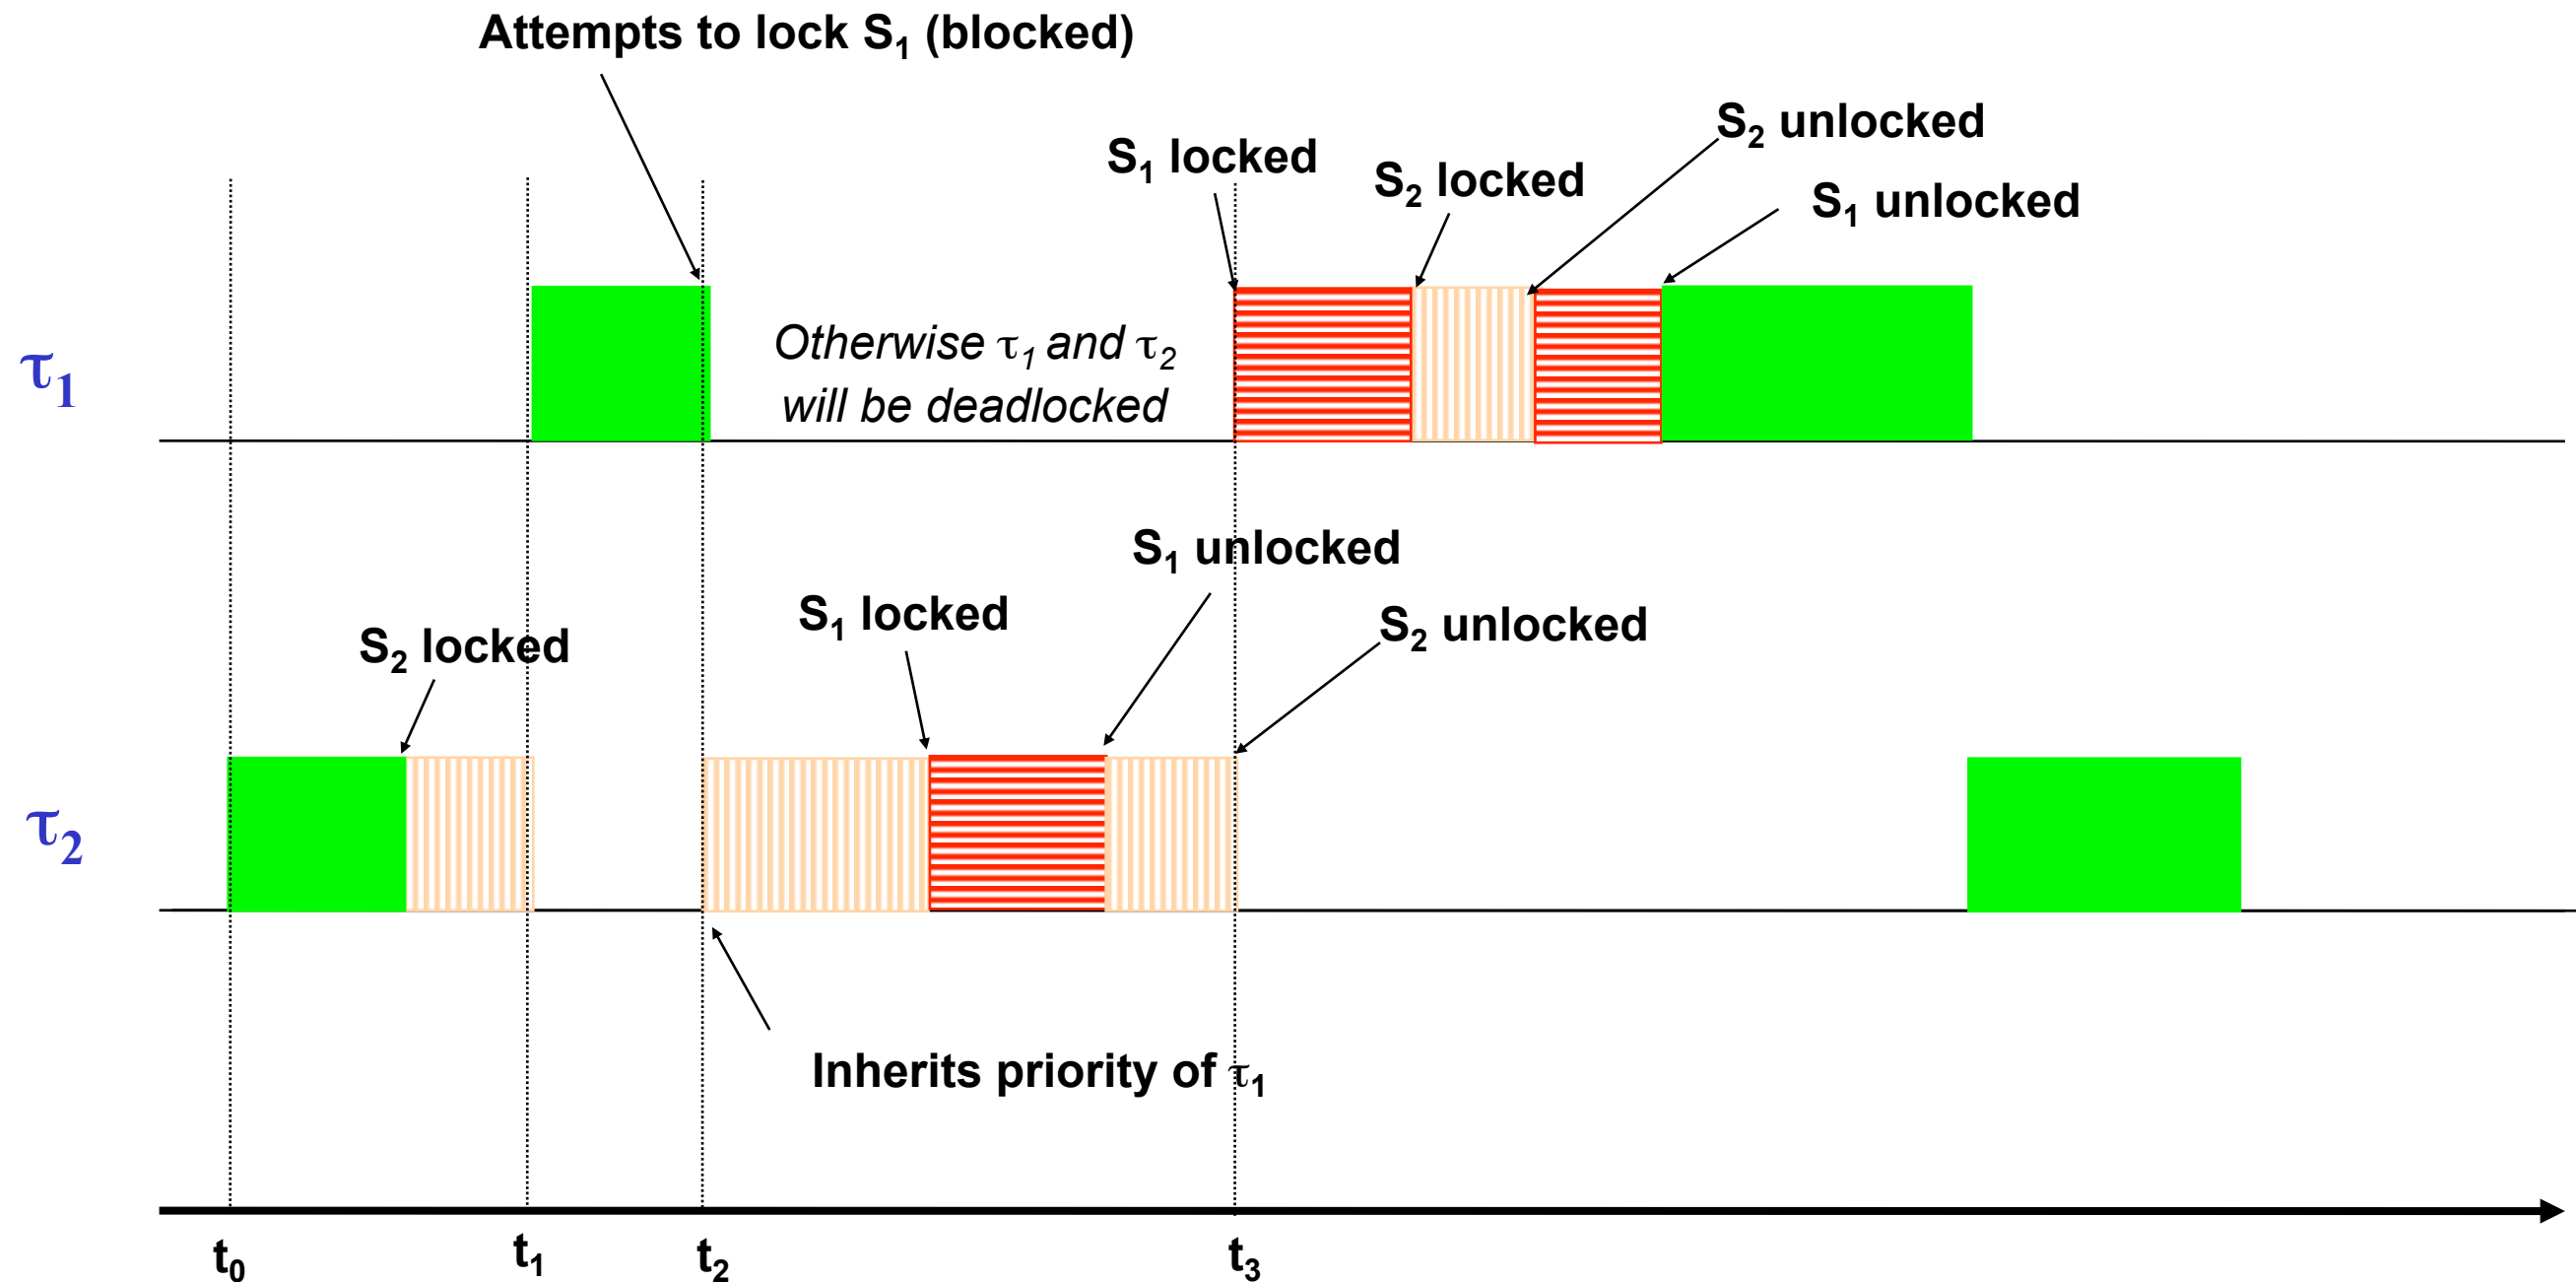
- Basic idea:
 - ▶ Priority ceiling of a binary semaphore S is the highest priority of all tasks that may lock S
 - ▶ When a task τ attempts to lock a semaphore, it will be blocked unless its priority is $>$ than the priority ceiling of all semaphores currently locked by tasks other than τ
 - ▶ If task τ is unable to enter its critical section for this reason, the task that holds the lock on its semaphore with the highest priority ceiling is
 - Said to be blocking τ
 - Hence, inherits the priority of τ

Example of Priority Ceiling Protocol

51

- Two tasks τ_1 and τ_2 with two shared data structures protected by binary semaphores S_1 and S_2 .
 - τ_1 : { ... Lock(S_1) ... Lock(S_2) ... Unlock (S_2) ... Unlock (S_1) ... }
 - τ_2 : { ... Lock(S_2) ... Lock(S_1) ... Unlock (S_1) ... Unlock (S_2) ... }
- Assume τ_1 has higher priority than τ_2
- Note: priority ceilings of both S_1 and S_2 = priority of τ_1

Example of Priority Ceiling Protocol



Priority Ceiling Protocols (contd.)


53

- Two forms
 - Original ceiling priority protocol (OCP)
 - Immediate ceiling priority protocol (ICPP)
- On a single processor system
 - A high priority process can be blocked at most once during its execution by lower priority processes
 - Deadlocks are prevented
 - Transitive blocking is prevented
 - Mutual exclusive access to resources is ensured (by the protocol itself)

- Each process has a static default priority assigned (perhaps by the deadline monotonic scheme)
- Each resource has a static ceiling value defined, this is the maximum priority of the processes that use it
- A process has a dynamic priority that is the maximum of its own static priority and the ceiling values of any resources it has locked.

Example of ICPP

55




- Each process has a static default priority assigned (perhaps by the deadline monotonic scheme)
- Each resource has a static ceiling value defined, this is the maximum priority of the processes that use it
- A process has a dynamic priority that is the maximum of its own static priority and any it inherits due to it blocking higher priority processes
- A process can only lock a resource if its dynamic priority is higher than the ceiling of any currently locked resource (excluding any that it has already locked itself).

$$B_i = \max_{k=1}^K usage(k, i) CS(k)$$

Example of OCPP

57



OCPP vs. ICPP

58

- Worst case behavior identical from a scheduling point of view
- ICPP is easier to implement than the original (OCPP) as blocking relationships need not be monitored
- ICPP leads to less context switches as blocking is prior to first execution
- ICPP requires more priority movements as this happens with all resource usages; OCPP only changes priority if an actual block has occurred.

Schedulability Impact of Task Synchronization

59

- Let B_i be the duration in which τ_i is blocked by lower priority tasks
- The effect of this blocking can be modeled as if τ_i 's utilization were increased by an amount B_i/T_i
- The effect of having a deadline D_i before the end of the period T_i can also be modeled as if the task were blocked for $E_i=(T_i-D_i)$ by lower priority tasks
 - As if utilization increased by E_i/T_i
- Theorem: A set of n periodic tasks scheduled by RM algorithm will always meet its deadlines if:

$$i, 1 \leq i \leq n, \frac{C_1}{T_1} + \frac{C_2}{T_2} + \dots + \frac{C_i + B_i + E_i}{T_i} \leq i(2^{1/i} - 1)$$

Arbitrary Deadlines

60

- Case when deadline $D_i < T_i$ is easy...
- Case when deadline $D_i > T_i$ is much harder
 - ▶ Multiple iterations of the same task may be alive simultaneously
 - ▶ May have to check multiple task initiations to obtain the worst case response time
- Example: consider two tasks
 - ▶ Task 1: $C_1 = 28, T_1 = 80$
 - ▶ Task 2: $C_2 = 71, T_2 = 110$
 - ▶ Assume all deadlines to be ∞

Arbitrary Deadlines (contd.)

61

- Response time for task 2:

Task 1: $C_1 = 28$, $T_1 = 80$

Task 2: $C_2 = 71$, $T_2 = 110$

- Initiation	Completion time	Response time
- 0	127	127
- 110	226	116
- 220	353	133
- 330	452	122
- 440	551	111
- 550	678	128
- 660	777	117
- 770	876	106

- Response time is maximum not for the first initiation of the task!
 - ▶ Not sufficient to consider just the first iteration
 - ▶ Theorem 1 (critical instant definition) no longer holds

Schedulability for Arbitrary Deadlines

62

- Analysis for situations where D_i (and hence potentially R_i) can be greater than T_i

$$w_i^{n+1}(q) = B_i + (q+1)C_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n(q)}{T_j} \right\rceil C_j$$

- The number of releases that need to be considered is bounded by the lowest value of $q = 0, 1, 2, \dots$ for which the following relation is true:

$$R_i(q) = w_i^n(q) - qT_i$$

$$R_i(q) \leq T_i$$

Arbitrary Deadlines (contd.)

63

- The worst-case response time is then the maximum value found for each q :

$$R_i = \max_{q=0,1,2,\dots} R_i(q)$$

- Note: for $D \leq T$, the relation $R_i(q) \leq T_i$ is true for $q=0$ if the task can be scheduled, in which case the analysis simplifies to original
 - If any $R > D$, the task is not schedulable

Dynamic Priority Scheduling

64

- With dynamic-priority scheduling, priorities are assigned to individual instances (jobs) of a task
- One of the most used algorithms of this class is EDF, or Earliest Deadline First priority assignment
 - Priorities assigned to tasks are inversely proportional to absolute deadlines of active jobs
 - It is optimal among all preemptive scheduling algorithms
 - If there exists a feasible schedule, then schedule given by EDF is also feasible
- Another optimal algorithm: Least Laxity First (LLF)
 - Assigns processor to the active task with smallest laxity
 - Larger overhead than EDF due to higher number of context switches
 - Less studied than EDF due to this reason

Preemptive Earliest Deadline First

65

- Processor executes the task whose absolute deadline is the earliest
- Priorities change with the closeness of a task to its absolute deadline
- Example:

Task	Arrival Time	Execution Time	Absolute Deadline
T1	0	10	30
T2	4	3	10
T3	5	10	25

EDF Schedulability

66

- Shown to be optimal for single processor
 - If EDF cannot schedule a task set on a single processor, then no other scheduling algorithm can
- Simple schedulability test if tasks are periodic, and have relative deadlines greater than or equal to their time periods
 - If the total utilization U of the task set is no greater than 1, the task set can be feasibly scheduled by EDF on a single processor
- If the relative deadlines are less than the time periods, there is no simple schedulability test
 - One will have to develop a schedule using EDF to see whether all the deadlines are met over a given time interval

EDF Schedulability ($D == T$ case)

67

- Theorem: Suppose we have a set of n periodic tasks (T_i, C_i, D_i) , each of whose relative deadline D_i equals its period T_i . The tasks can be feasibly scheduled by EDF on a single processor iff:

$$C_1/T_1 + C_2/T_2 + \dots + C_n/T_n \leq 1$$

Summary

68

- Real-Time Scheduling: Orchestrating the execution of multiple tasks (processes) so that timing constraints of tasks are satisfied
- Preemptive, priority based scheduling is the most commonly used approach
 - Static priority assignment: Rate Monotonic, Deadline Monotonic
 - Dynamic priority assignment: Earliest Deadline First
- Key results on schedulability exist for both RM and EDF. These results allow us to model and analyze an embedded software system to understand its feasibility before actually building it
- Several real-world effects (shared resources leading to priority inversion, etc.) can also be accounted for